

Scheduling

Motivation

- especially for behaviors
 - pseudo-parallelism
 - i.e., processes share one (few) CPU(s) to emulate concurrent execution
- but also relevant for “higher” layers
 - modeling, planning, reasoning, etc. can profit from concurrency
- note: due to use of “normal” computers
 - Grey Walter’s tortoises (late 40s - 50s)
 - analog circuits (radio vacuum tubes)
 - differential drive, light sensors, bumper
 - homing in on white light => charging station
 - inherently parallel computations

Grey Walter's Tortoises



Scheduling

- purpose
 - making use of a scarce resource
 - typically processing power
 - but also harddisk, printer, etc. access
- time sharing aka time slicing
 - **time slot** aka **quantum**: constrained time interval for execution
 - **context switch**: change from one executed process to next one

Terminology

- task / process
 - "large", application
 - "independent", especially own address space
- thread
 - "small", part of a program (i.e., task/process)
 - shared address space
- no fixed terminology
 - depending e.g. on type of OS
- here:
 - process as term for the general category
 - task and thread as subclasses

Important Process Constraints

- deadline
 - a moment in time until when
 - the execution of the process must be finished
- ready time
 - earliest possible moment in time
 - where the process may be started
- period
 - minimal frequency at which a time slot
 - for a cyclic process must be provided
 - i.e., regular deadline

Realtime Processing

- hard real-time
 - fullfill all constraints at all times
 - e.g., control of critical components
 - often computed offline
- soft real-time
 - best effort
 - e.g., multimedia streams
- implementation of hard realtime
 - determination of max. runtime of a process
 - problems with normal OS/hardware
 - especially pipelining, cacheing, interrupts

Two major approaches

- cooperative scheduling
 - running process stops by itself
 - and invokes the scheduler
- preemptive scheduling
 - scheduler stops (preempts) running process
 - higher demands on the context switch

Finding an optimal schedule

- naive algorithm
 - test all possibilities
 - exponential runtime
- is a better algorithm possible?
 - n non-preemptive processes
 - with ready-times, execution times, deadline
 - searching optimal schedule is NP-complete
 - can be reduced to bin-packing

$$n! \approx \left(\frac{n}{e}\right)^n \cdot \sqrt{2\pi n}$$

All processes created equal, or not?

- round robin scheduling
 - most simple and very commonly used
 - FIFO principle
- but
 - not all processes of same "importance"
 - need for "preferences"
 - => use of priorities

B-Scheduling

Behaviors: Need of Scheduling

- several behaviors “active” at the same time
 - pseudo-parallelism
 - need to decide which process runs next
- simple solution
 - run one after the other
 - i.e., round robin scheduling

Problems RR behavior scheduling

- many behaviors at different time scales
 - e.g., motor control vs battery monitor
 - proper scheduling with priorities needed
- behaviors are used for control
 - not only deadlines
 - but also minimization of jitter,
 - i.e., variations in the execution period

Handling different Time-Scales: Exponential Effect Priorities

exponential effect:

priority incremented

=> frequency halved

- process $P[i]$
- priority $prio[i]$
 - $prio[x] == 0$:
maximum frequency
 - $prio[x] == n$:
1/2 frequency of $n-1$

example

priority	frequency
0	1024 Hz
1	512 Hz
2	256 Hz
3	128 Hz
4	64 Hz
...	...
10	1 Hz
11	every 2 sec
...	...
16	~1 minute
...	...
22	~1.1 hour