


Problem Solving as Search


Problem-Solving

- Goal Formulation
 - set of one or more (desirable) world states
 - e.g. checkmate in chess
- Problem formulation
 - what actions and states to consider
 - given a goal and an initial state
- Search for solution
 - given the problem, search for a solution,
 - i.e., a sequence of actions to achieve the goal starting from the initial state
- Execution of the solution

Example: Path Finding problem

- Formulate goal:
 - be in Bucharest (Romania)
- Formulate problem:
 - initial state: be in Arad
 - action: drive between pair of connected cities (direct road)
 - state: be in a city (20 world states)
- Find solution:
 - sequence of cities leading from start to goal state, e.g., Arad, Sibiu, Fagaras, Bucharest
 - Execution
 - drive from Arad to Bucharest according to the solution






Problem-Solving


- hence conceptually very easy
- but: where's the catch?
(or why haven't the robot's enslaved us yet?)

the curse of combinatorial explosion
(or our savior 😊)

Example: Blocks World




(a) "Pick up a big red block."



(b) "Find a block which is taller than the one you are holding and put it into the box."

How many
different possible
world states?

> 4.5 billion states
even in this
toy domain!



(c) "Will you please stack up both of the red blocks and either a green cube or a pyramid?"

- a) Tens?
- b) Hundreds?
- c) Thousands?
- d) Millions?
- e) Billions?
- f) Trillions?

Common trick: state evaluation functions aka “heuristics”

- provide guidance w.r.t. what action to take next
- avoids exploring the whole search space
- e.g.,
 - consider all neighboring states, which are reachable via some action
 - select the action that leads to the state with the highest utility (evaluation value)
 - i.e., greedy approach to action selection

Search Methods

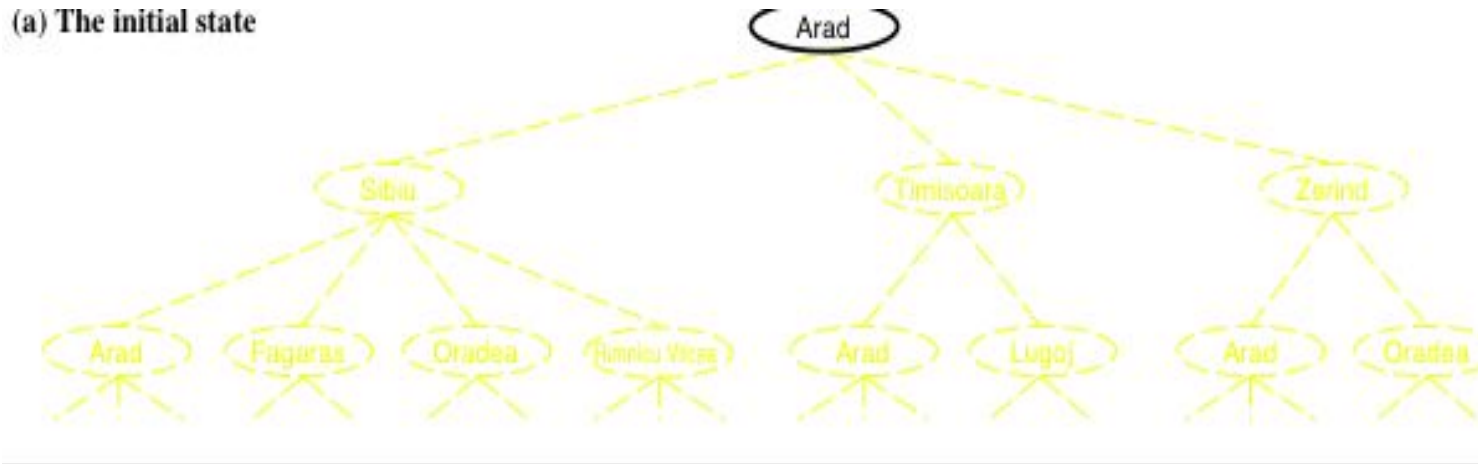
Basic search algorithms

Problem-Solving

- search the state space
- e.g., through *explicit tree generation*
 - ROOT= initial state
 - nodes and leafs generated through successor function

Simple tree search example

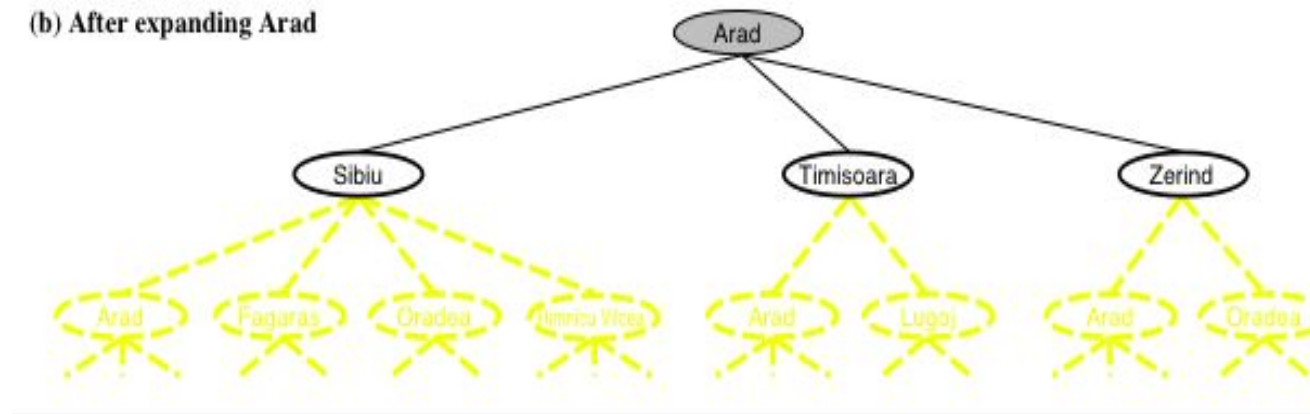
(a) The initial state



```
function TREE-SEARCH(problem, strategy) return a solution or failure
  Initialize search tree to the initial state of the problem
  do
    if no candidates for expansion then return failure
    choose leaf node for expansion according to strategy
    if node contains goal state then return solution
    else expand the node and add resulting nodes to the search tree
  enddo
```


Simple tree search example

(b) After expanding Arad



```
function TREE-SEARCH(problem, strategy) return a solution or failure
  Initialize search tree to the initial state of the problem
  do
    if no candidates for expansion then return failure
    choose leaf node for expansion according to strategy
    if node contains goal state then return solution
    else expand the node and add resulting nodes to the search tree
  enddo
```

Simple tree search example



```
function TREE-SEARCH(problem, strategy) return a solution or failure
  Initialize search tree to the initial state of the problem
  do
    if no candidates for expansion then return failure
    choose leaf node for expansion according to strategy
    if node contains goal state then return solution
    else expand the node and add resulting nodes to the search tree
  enddo
```

Tree search algorithm

```
function TREE-SEARCH(problem, fringe) return a solution or failure
  fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if EMPTY?(fringe) then return failure
    node ← REMOVE-FIRST(fringe)
    if GOAL-TEST[problem] applied to STATE[node] succeeds
      then return SOLUTION(node)
    fringe ← INSERT-ALL(EXPAND(node, problem), fringe)
```

FRINGE (aka ***frontier***)

= contains generated nodes which are not yet expanded

Tree search algorithm (2)

```
function EXPAND(node,problem) return a set of nodes
    successors  $\leftarrow$  the empty set
    for each  $\langle$ action, result $\rangle$  in SUCCESSOR-FN[problem](STATE[node]) do
        s  $\leftarrow$  a new NODE
        STATE[s]  $\leftarrow$  result
        PARENT-NODE[s]  $\leftarrow$  node
        ACTION[s]  $\leftarrow$  action
        PATH-COST[s]  $\leftarrow$  PATH-COST[node] + STEP-COST(node, action, s)
        DEPTH[s]  $\leftarrow$  DEPTH[node]+1
        add s to successors
    return successors
```

Search strategies

search strategy: picking the order of node expansion

- problem-solving performance:
 - **Completeness**: *always find a solution if one exists?*
 - **Optimality**: *always find the least-cost solution?*
 - **Time Complexity**: *#nodes generated/expanded?*
 - **Space Complexity**: *#nodes in memory during search?*
- measured in terms of problem difficulty:
 - *b : maximum branching factor of the search tree*
 - *d : depth of the least-cost solution*
 - *m : maximum depth of the state space (may be ∞)*

Uninformed search strategies


aka blind search

- use only information in problem definition
- example categories of expansion
 - Breadth-first search (BFS)
 - Depth-first search (DFS)
 - Iterative deepening search
 - Bidirectional search

(informed search: strategies can determine whether one state is better than another)


BFS

- Expand *shallowest* unexpanded node
- Implementation: *fringe* is a **FIFO** queue (First In, First Out)




BFS

- Expand *shallowest* unexpanded node
- Implementation: *fringe* is a FIFO queue




BFS

- Expand *shallowest* unexpanded node
- Implementation: *fringe* is a FIFO queue



BFS

- Expand *shallowest* unexpanded node
- Implementation: *fringe* is a FIFO queue



BFS evaluation

Completeness:

- Does it always find a solution if one exists?
- YES
 - If shallowest goal node is at some finite depth d
 - Condition: If b is finite, i.e., maximum number of successor nodes is finite

BFS evaluation

- Completeness:
 - YES (if b is finite)
- Time complexity:
 - assume branching factor b (all nodes have b successors)
 - i.e., root has b successors, each node at the next level has again b successors (total b^2), ...
 - solution is at depth d
 - worst case: expand all but the last node at depth d
 - total #nodes expanded:


$$b + b^2 + b^3 + \dots + b^d + (b^{d+1} - b) = O(b^{d+1})$$

BFS evaluation

- Completeness:
 - YES (if b is finite)
- Time complexity:
 - Total numb. of nodes generated: $O(b^{d+1})$
- Space complexity:
 - like time if each node is retained in memory
- Optimality:
 - *Does it always find the least-cost solution?*
 - In general YES (unless actions have non-uniform cost, i.e., deeper states can be cheaper)


DFS

- Expand *deepest* unexpanded node
- Implementation: *fringe* is a LIFO queue (i.e., stack)




DFS

- Expand *deepest* unexpanded node
- Implementation: *fringe* is a LIFO stack




DFS

- Expand *deepest* unexpanded node
- Implementation: *fringe* is a LIFO stack




DFS

- Expand *deepest* unexpanded node
- Implementation: *fringe* is a LIFO stack




DFS

- Expand *deepest* unexpanded node
- Implementation: *fringe* is a LIFO stack




DFS

- Expand *deepest* unexpanded node
- Implementation: *fringe* is a LIFO stack




DFS

- Expand *deepest* unexpanded node
- Implementation: *fringe* is a LIFO stack




DFS

- Expand *deepest* unexpanded node
- Implementation: *fringe* is a LIFO stack




DFS

- Expand *deepest* unexpanded node
- Implementation: *fringe* is a LIFO stack




DFS

- Expand *deepest* unexpanded node
- Implementation: *fringe* is a LIFO stack




DFS

- Expand *deepest* unexpanded node
- Implementation: *fringe* is a LIFO stack



DFS

- Expand *deepest* unexpanded node
- Implementation: *fringe* is a LIFO stack



DFS evaluation

Branching factor – b

Maximum depth – m

- Completeness
 - NO unless search space is finite
 - may go infinitely into one branch without solution
- Time complexity $O(b^m)$
 - terrible if m is much larger than the depth d of the optimal solution
 - but **if there are many solutions**, then it is **typically faster than BFS**

DFS evaluation

- Completeness:
 - NO unless search space is finite
- Time complexity: $O(b^m)$
- Space complexity: $O(bm + 1)$
 - backtracking search uses even less memory
 - one successor instead of all b
- Optimality:
 - NO unless search space is finite

Iterative deepening search

function ITERATIVE_DEEPENING_SEARCH(*problem*) **return** a solution or failure

inputs: *problem*

for *depth* $\leftarrow 0$ to ∞ **do**
 result \leftarrow DEPTH-LIMITED_SEARCH(*problem*, *depth*)
 if *result* \neq cutoff **then return** *result*

- DFS with increasing limits
- combines benefits of DFS and BFS

ID-search, example

- Limit=0




ID-search, example

- Limit=1




ID-search, example

- Limit=2



ID-search, example


- Limit=3



IDS evaluation

- Completeness:
 - YES (no infinite paths)
- Time complexity: $O(b^d)$
- Space complexity: $O(bd)$
- Optimality:
 - YES (if step cost is uniform)

Bidirectional search




- Two simultaneous searches from start and goal
- Check if node belongs to other fringe before expansion
- Motivation: $b^{d/2} + b^{d/2} \neq b^d$
- Complete and optimal if both searches are BFS
- Predecessor of each node should be efficiently computable

Summary of algorithms

Criterion	Breadth-First	Depth-First	Depth-limited	Iterative deepening	Bidirectional search
Complete?	YES*	NO	YES, if $l \geq d$	YES	YES*
Time	b^{d+1}	b^m	b^l	b^d	$b^{d/2}$
Space	b^{d+1}	bm	bl	bd	$b^{d/2}$
Optimal?	YES*	NO	NO	YES	YES

Repeated states

Failure to detect repeated states can turn a solvable problem into unsolvable ones.



includes all problems with
reversible transitions/actions!!!

Solution: Graph Search

- use an “explored” set
to store already visited states/nodes
- expand only nodes from the frontier
that are not in the explored set