

Computer Vision

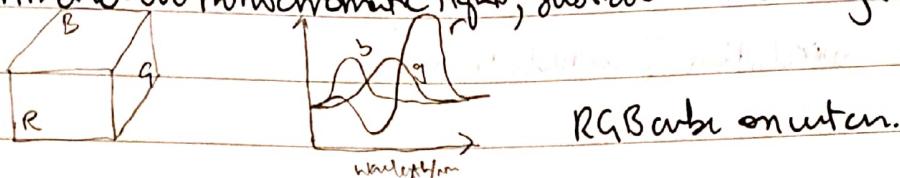
(1) Light sum or difference. How would that change?

Add: Colors add (combine) by adding color spectra (light adds to existing black).

Subtract: Colors combine by multiplying color spectra. (pigments remove color from incident white light).

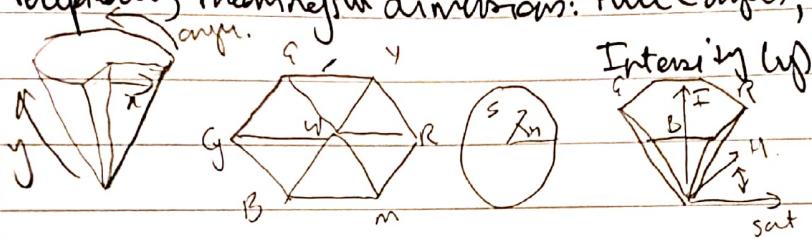
(2) Difference between RGB & HSV:

RGB: primaries on nonchromatic light, subtractive matching for some wavelength.



RGB cube on white.

HSV: perceptually meaningful dimensions: hue (angle), saturation (n)



(3) Different Binary & Grayscale:

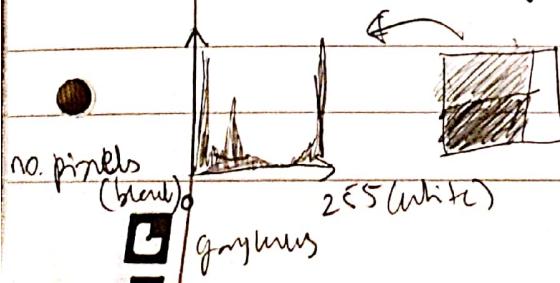
Binary: represents black (1) & white (0) in a [pxx] matrix.

Grayscale: represents 0 → 255 (white) [0 black]. [pxx] matrix
(intensity)

(4) Creating a histogram (numerical):

Histogram: provides the frequency of the brightness (intensity value) in an img.

[distribution of gray levels in an image] & [how frequently each gray level occurs]



(5) Define a filter:

Filters form a new image whose pixels values are transformed from the original pixel value.

(6) Apply a filter:

[look at examples & apply formulae]

(7) Difference correlation & convolution:

corr →

correlation: ($\ast \ast$) is a measure of similarity between 2 signals.

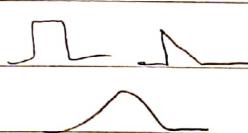
flip 180° → convolution: (\ast) is a linear operator on a signal which modifies it.

(used for blurring, sharpening / edge detection)

con



cor



minimum 2 signals must match

the amount of overlap
of convolution as
shifted one with another.

(8) How do sharpen an image?

original + details = sharpen.

$$\# \begin{bmatrix} 0 & a & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix} + a \begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix} - \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

$$\text{original } \# \begin{bmatrix} 0 & 0 & 0 \\ 0 & 0 & 0 \\ 0 & 0 & 0 \end{bmatrix} = \boxed{\quad} \xrightarrow{\text{(blurred with box filter)}} \text{shifted by 1 pixel.}$$

(A) Derivations: (forward) (backward) (center).

Back: $\frac{df}{dn} = f(m) - f(m-1) = f'(m) \rightarrow [0 \ 1 \ -1]$

forward: $\frac{df}{dn} = f(m) - f(m+1) = f'(m) \rightarrow [-1 \ 1 \ 0]$

center: $\frac{df}{dn} = f(m+1) - f(m-1) = f'(m) \rightarrow [1 \ 0 \ -1]$

(B) Edge derivatives

Gradient vector: $\nabla f(n, y): \begin{bmatrix} \frac{\partial f(n, y)}{\partial x} \\ \frac{\partial f(n, y)}{\partial y} \end{bmatrix} = \begin{bmatrix} f'_x \\ f'_y \end{bmatrix}$

gradient magnitude: $|\nabla f(n, y)| = \sqrt{f_x'^2 + f_y'^2}$

$$\frac{1}{3} \begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ 1 & 1 & 1 \end{bmatrix} \rightarrow \text{gradient direction}$$

gradient direction: $\theta = \tan^{-1}\left(\frac{f'_y}{f'_x}\right)$

$$\frac{1}{3} \begin{bmatrix} 0 & 0 & 0 \\ -1 & 1 & -1 \\ -1 & -1 & 1 \end{bmatrix} \rightarrow \text{gradient magnitude}$$

(C) Noise attacking derivatives.

Noise loses the edge, because smooth change occurring. \therefore

Results in pixels look very different from neighbor.

(Smoothing helps by forcing neighboring pixels to look similar)

mean: $\begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix} \xrightarrow{[1 \ 1 \ 1]} \text{opinion: } \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \xrightarrow{[1 \ 2 \ 1]} \text{(smoothing \& derivative)}$

\rightarrow blur image \rightarrow reduce contrast (smooth surface \rightarrow low contrast)

(D) Best filter for which noise? (mean \rightarrow low filter)

(E) Salt & pepper noise \rightarrow MEDIAN FILTER.

(F) noise reduction \rightarrow super resolution.

(G) mended up \rightarrow in-painting.

(13) Effect of smoothing (size window)

Reduces noise but also blurs image \rightarrow size of window inc. \therefore blurring.

(14) Criterion for edge detector:

(1) Good detection: minimize probability of false positives
(spurious edges by noise) (missing real edges)

(2) Good localization: edges detected must be as close as possible
to true edges

(3) Single response: return one point only for each true edge point.
(minimize no. of local minima around true edge)

(15) Canny thresholds

Canny uses 2 thresholds: (upper & lower), If a pixel higher than
upper threshold than it is an edge pixel: if lower than lower threshold
then rejected as an edge.

(16) Size of gaussians

σ (Gaussian kernel spread / size) \rightarrow σ large then large scale edges
detected, small σ detects fine features.

(17) Hough vs Ransac for line detection.

Hough can detect lines, circles and other structures only if the parametric
form is known (robust detection under noise & partial occlusion)

Pros (simple, easy implementation, missing '3' connected islands) adopt many forms

Cons (Complex computation for objects with many parameters, looks for one single
type of object, can be fooled by apparent lines, length & position of line
segment cannot be determined, collinear line segments cannot be
separated).

RANSAC (Random sample consensus):

divides data into inliers & outliers and yields estimate computed from minimal set of inliers. estimates best line with most inliers & least outliers

Pros: (General method suited for a wide range of model fitting problems, easy to implement & easy to calculate failure rate). $(1 - \alpha^n)^k$ (n -points to fit) ($\alpha = \text{fraction of inlier pts}$) ($k = \text{number of iterations}$) ($\alpha = \text{target failure rate}$)

Cons: (only handles moderate % of outliers without tuning up, may need more iterations for high rate of outliers)

(18) Can they be used for things other than lines? why?

Through transform yes! if you give the parametric equations.

Because yes, it estimates the parameters of a model by random sampling.

(19) Derivations & terms

at a corner the image gradient has at least 2 dominant directions

Corners are repeatable & distinctive

-locality: easily recognize the point by looking through overall window.

(Localization)- shifting the window in any direction should give always close intensity

corner: $\sum I_n^2$ large (constant) $\sum I_{ij}^2$ low.

$$\text{(Cross difference)} E(u, v) = \sum_{ij} w_{ij} I_{(n+u, n+v)} [I_{(n+u, n+v)} - I_{(n, n)}]^2 \\ M = \begin{bmatrix} \sum I_n^2 & \sum I_n \cdot I_n \\ \sum I_n \cdot I_n & \sum I_{ij}^2 \end{bmatrix} = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} \text{ (symmetric.)}$$

$$\text{Covariance rotation invariant: } M = g(Co)^T \begin{bmatrix} I_n^2 & I_{nj} \\ I_{nj} & I_{jj} \end{bmatrix}$$

(20) Harris detector (properties) proc.

1. Translation invariant
 2. Rotation invariant
 3. Not scale invariant.
1. image derivatives
 2. second derivatives
 3. gaussian filter $g(\sigma)$
4. corner but (strong edge)
 - $\theta = \det[M(c_0, \sigma_0)] - \text{tr}[M(c_0, \sigma_0)]^2$
 - $= g(I_x^2)g(I_y^2) - [g(I_x)]^2$
 - $- d[g(I_x^2) + g(I_y^2)]$
- (5) Non max suppression.

(21) SIFT

Finds the local minimum using different gaussian in open² scale.

Captures key points at varying scales \rightarrow same interest points independently in

with scaled base.

- make blurred image associate with the key points scale
- to know rotation invariant, note the gradient directions & locations by keypoint orientation.
- create array of orientation histograms. (Gradient = 8 bins = 128 int.)
- it is invariant to rotation because worked with rotated derivative
- invariant to scale worked with scaled image in ROI
- compare each vector in A to vectors in B to find matching keypoints (nearest neighbor for high contrast).

Very large illumination gradients \rightarrow abelities ≤ 0.2 the normalize.
(invariant to illumination)

11)

(22) Affinitive cluster.

Clustering: group together similar data points & represent them with a single token.

- look at large amt of data, partition based compression / division,
- represent a large cluster with a token with the cluster number.

(Histogram feature, color, sift vector)

Significant regions.

Predict image in some clusters somehow

Agglomerative clustering: start each point as its own cluster & iteratively merge the closest cluster.

(1) every point is own cluster.

(2) find most similar pair \rightarrow merge

(3) repeat

[
Any distance, min distance, mindist.
(e.g. L1)
set threshold \rightarrow once above threshold
stop clustering, mindist between clusters
(longest chain)
Complete Link \rightarrow num distance points (right cluster).]

(23) k-means cluster

(1) Initialize cluster centers ($c_1, c_2 \dots c_n$).

→ set assignment for each point to cluster c_i .

(2) Assign each point to nearest cluster (compute δ^t) → iteration t.

$$\delta^t = \arg\min_{\delta} \frac{1}{N} \sum_j \sum_i \delta_{ij}^{t-1} (c_i^{t-1} - x_j)^2 \quad \text{[Euclidean distance used]}$$

(3)

(3) Compute c^t : update cluster centers as the mean of the points.

$$c^t = \arg\min_c \sum_j \sum_i \delta_{ij}^{t-1} (c_i - x_j)^2$$

(24) K-mean vs mean shift

K-mean:	finds cluster centers that minimize conditional variance.
P	simple and fast. most robust
C	susceptible to outliers more to local minima All clusters have same prior (deterministic is not adaptive)

mean shift: application independent

	model free, does not demand prior shift
P	single parameter (window size). finds variable no. of nodes, robust to outliers.
C	outperforms on midsize window size is not trivial computationally complex does not scale well with dimensions of feature space

(25) Pipeline for object recognition:

Apply a prediction function to a feature representation of the image
no yet disentangled: $f(\text{apple}) = \text{apple}$

Training images \rightarrow image features \rightarrow training \rightarrow learned classifier

Testing images \rightarrow image features \rightarrow learned classifier \rightarrow prediction

image features: color, shape, texture

(quantize RGB) (PCA) (filter banks)

(nearest neighbor)

(26) (boy or march)
BOW: why weight?

Stochastic BM uses extract features and build a dictionary of visual words.

[Given a query extract features & build histogram for each feature [and closest visual word].

Weight because some visual words are more discriminative than others
(the higher the rank of the document the word appears in the less useful it is).

(27) Camera model parameters:

intrinsic assumptions:

[unit aspect ratio]

optical center (0,0)

no skew

Extrinsic Assumptions:

[no rotation]

[constant 0,0,0]

$$\text{From: } P' = \begin{pmatrix} x \\ y \\ z \end{pmatrix}$$

$$P' = \begin{pmatrix} I & 0 & 0 \\ 0 & I & 0 \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} x \\ y \\ z \end{pmatrix} = k \begin{pmatrix} I & 0 \end{pmatrix} P$$

$$P' = k(I_0)P \rightarrow = \begin{pmatrix} I & 0 & 0 \\ 0 & I & 0 \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} u_0 \\ v_0 \\ 1 \end{pmatrix} \begin{pmatrix} x \\ y \\ z \end{pmatrix} \xrightarrow{\text{center.}}$$

$$\text{Square pixels: } \rightarrow \text{next: } P' = \begin{pmatrix} u_0 & v_0 & 0 \\ 0 & u_0 & v_0 \\ 0 & 0 & 1 \end{pmatrix}$$

$$\text{Skewness: } \begin{pmatrix} u_0 & v_0 & 0 \\ 0 & u_0 & v_0 \\ 0 & 0 & 1 \end{pmatrix}$$

$$\begin{pmatrix} u_0 & v_0 & 0 \\ 0 & u_0 & v_0 \\ 0 & 0 & 1 \end{pmatrix} \xrightarrow{\text{intrinsic params}} \begin{pmatrix} r_{11} & r_{12} & r_{13} & t_x \\ r_{21} & r_{22} & r_{23} & t_y \\ r_{31} & r_{32} & r_{33} & t_z \end{pmatrix} \xrightarrow{\text{extrinsic}}$$

$$\text{SIFT}$$

$$\text{GMMF}$$

(28) Dot projection matrix
5 6

intrinsic extrinsic.

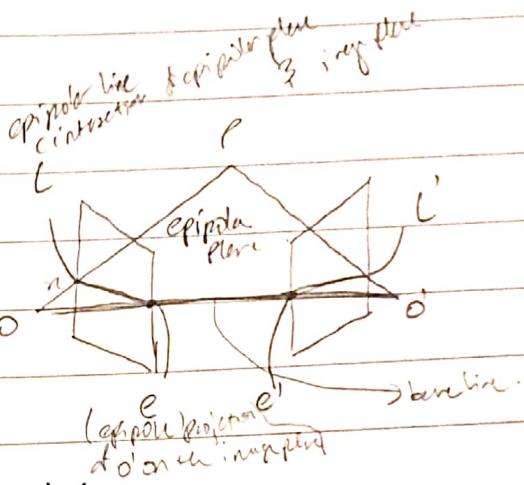
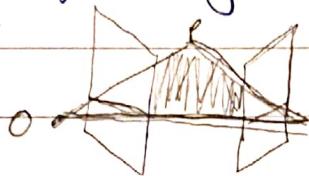
(29) Epipolar geometry:

Given a point in camera

$$BEx_n = l'$$

antibody in essential

matrix gives epipole line in 2nd view.



at $Ex_n = 0$ The point maps to a line l' in right image.

$n^T l = 0$ The baseline contains the $l \cap l'$. c \cap c'?

\hookrightarrow epipole An epipole e is a projection of d on the image plane

$E = R[t_x]$ all epipole lines in an image intersect at the epipole.

(30) 8-point alg

(1) normalize points

G) construct mag matrix A for matched pts. $x_m^T F x_m = 0$

(2) find SVD of A.

(3) entries of f are elements of column V corresponds to least singular value.

(4) extract rank 2 from F

(5) unnormalize $F \leftarrow (SVD)$

\therefore geometric?

$$\begin{aligned} & \begin{bmatrix} x_m & y_m & 1 \end{bmatrix} \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x_m \\ y_m \\ 1 \end{bmatrix} = 0 \\ & \begin{bmatrix} x_1 & x_1 & 1 & x_2 & y_2 & 1 & x_3 & y_3 & 1 & x_4 & y_4 & 1 \end{bmatrix} \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x_1 & y_1 & 1 \\ x_2 & y_2 & 1 \\ x_3 & y_3 & 1 \\ x_4 & y_4 & 1 \end{bmatrix} = 0 \\ & \begin{bmatrix} x_{m1} & x_{m2} & x_{m3} & x_{m4} & y_{m1} & y_{m2} & y_{m3} & y_{m4} & 1 & x_{m1} & y_{m1} & 1 \\ x_{m1}^2 & x_{m2}^2 & x_{m3}^2 & x_{m4}^2 & y_{m1}^2 & y_{m2}^2 & y_{m3}^2 & y_{m4}^2 & 1 & x_{m1} y_{m1} & 1 \end{bmatrix} \begin{bmatrix} f_1 \\ \vdots \\ f_9 \end{bmatrix} = 0 \end{aligned}$$

(32) Steps stereo rectification:

Reproject image plane onto a common plane parallel to the line between camera centers

1) Rotate the right camera by R . [compute E w.r.t R]

2) Rotate (wrt left) the left camera so that the epipole is at ∞ .

3) Rotate (wrt left) right camera so epipole @ ∞ .

4) adjust scale.

$$T_1 = U_3 \quad T_2 = -U_3 \quad E = [R \ I \ T]$$
$$R_1 = U_1 W V^T \quad R_2 = U_2 W^T V^T \quad (\text{choose } d=1)$$

$$\begin{bmatrix} 0 & -1 & 0 \\ 1 & 0 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

rotate both, rotate right by $R \rightarrow$ rotate both steps by R instead.

Scale Both by H !

(33) Why rectification?

to do stereo matching etc...

(34) Assumptions for optical flow:

Planar image motion of an object leads from spatio-temporal image brightness variations (optical flow) \rightarrow measure local features we want to track optical flow algorithms can help track those.

Assume: the brightness of any object is constant over time

- nearby points in the image plane move in a similar manner (velocity smoothness constraint)

(35) derivation?