



JACOBS
UNIVERSITY

Query Processing and Optimization

Jennifer Widom

Ramakrishnan/Gehrke Chapters 10, 12

Steps in Database Query Processing

Parser – Checker - Views - Logical plan – Optim1 - Physical plan – Optim2 - Execution

Query string

→ **Parser** →

Query tree

→ **Checker** →

Valid query tree

→ **View expander** →

Valid tree w/o views

→ **Logical query plan generator** →

Logical query plan

→ **Query rewriter (heuristic)** →

Better logical plan

→ **Physical query plan generator (cost-based)**

Selected physical plan

→ **Code generator** →

Executable code

→ **Execution engine**

Running Example

Parser – Checker - Views - Logical plan – Optim1 - Physical plan – Optim2 - Execution

- Tables (what are the keys?):

Student(ID, Name, Major)

Course(Num, Dept)

Taking(ID, Num)

- Query to find all EE students taking at least one CS course:

```
SELECT Name
FROM   Student, Course, Taking
WHERE  Taking.ID = Student.ID
      AND Taking.Num = Course.Num
      AND Major = 'EE'
      AND Dept = 'CS'
```

π

\times

\bowtie

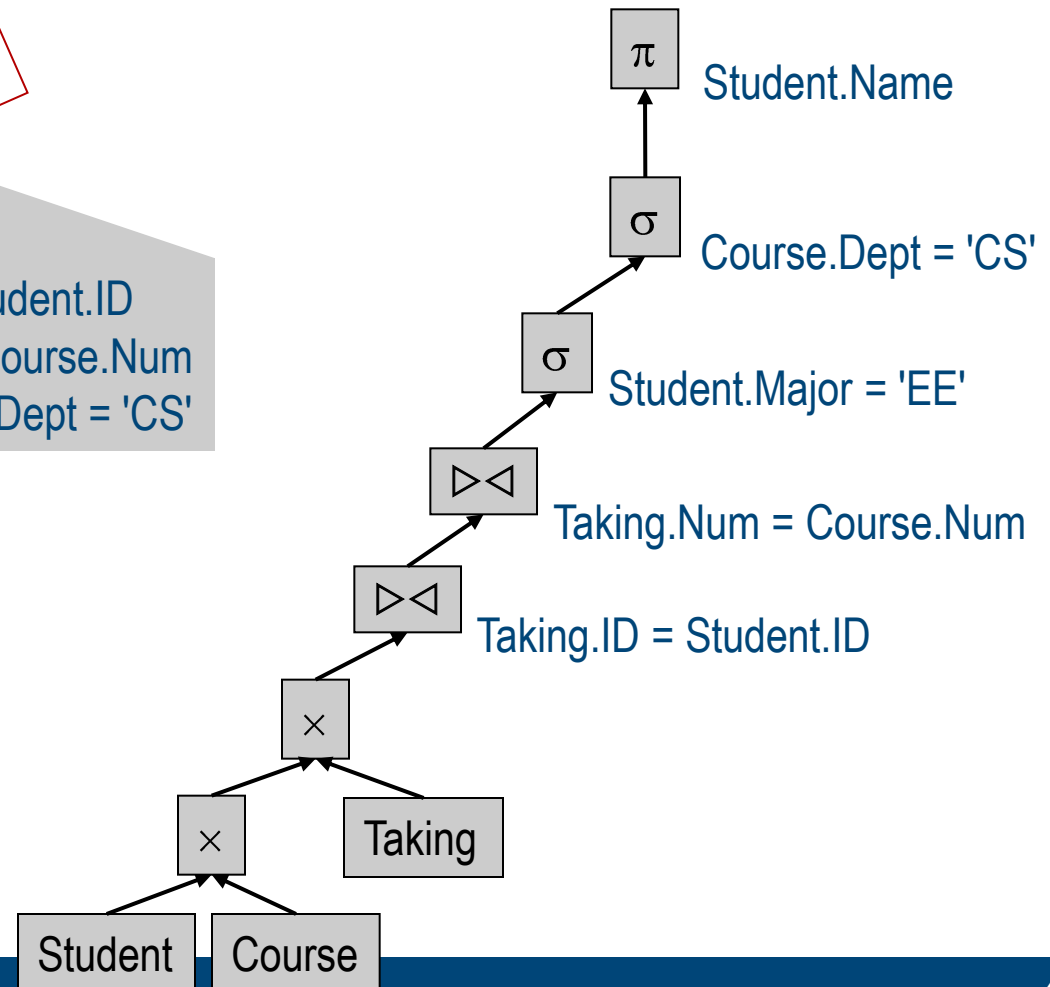
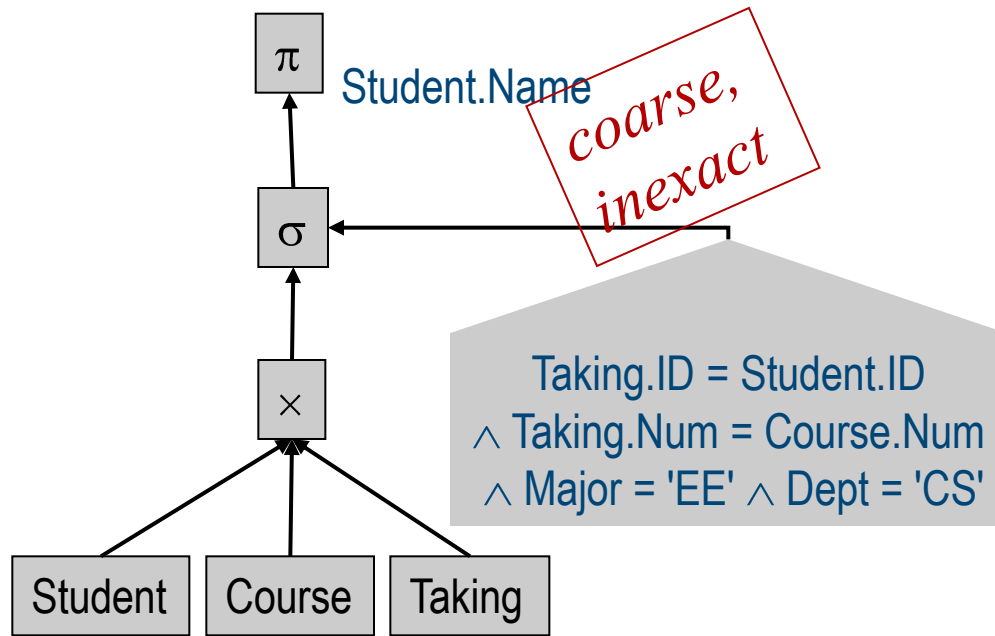
\bowtie

σ

σ

Logical Query Plan

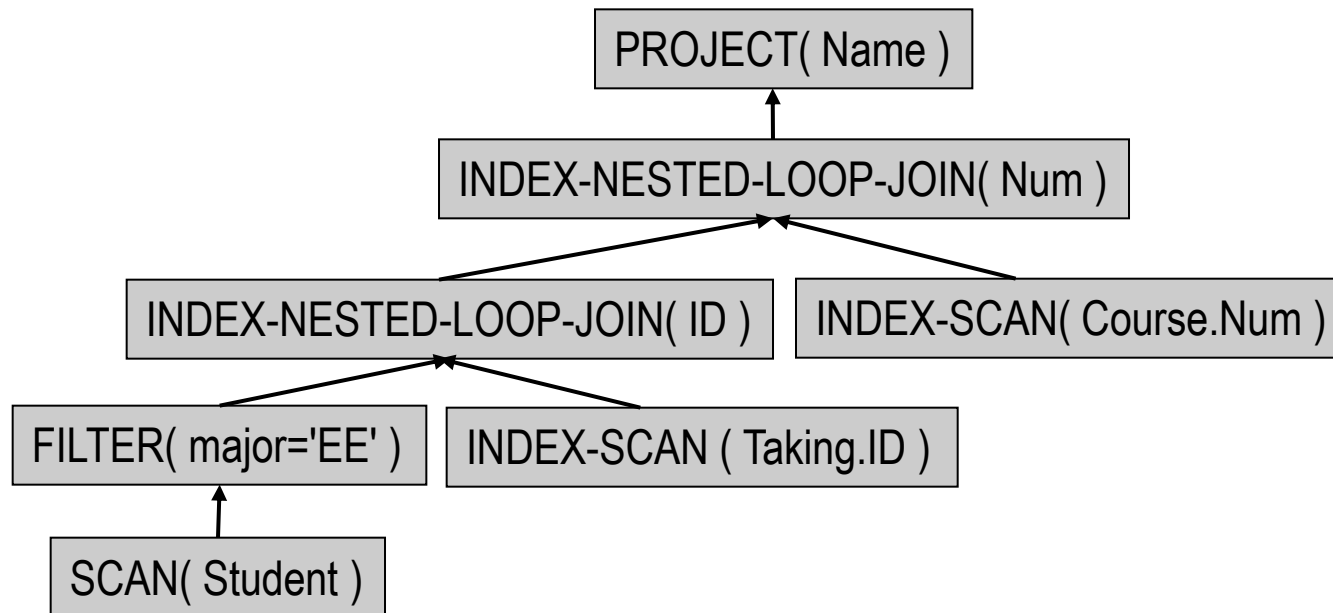
Parser – Checker - Views - **Logical plan** – Optim1 - Physical plan – Optim2 - Execution



```
SELECT Name
FROM   Student, Course, Taking
WHERE  Taking.ID = Student.ID
      AND Taking.Num = Course.Num
      AND Major = 'EE'
      AND Dept = 'CS'
```

Physical Query Plan

Parser – Checker - Views - Logical plan – Optim1 - **Physical plan** – Optim2 - Execution



```
SELECT Name
FROM Student, Course, Taking
WHERE Taking.ID = Student.ID
      AND Taking.Num = Course.Num
      AND Major = 'EE'
      AND Dept = 'CS'
```

*one of manyManyMany possible plans,
assumes particular index situation.*

Sample Operator: Nested Loop Join

Parser – Checker - Views - Logical plan – Optim1 - **Physical plan** – Optim2 - Execution

- Consider this equi-join query:

```
SELECT *  
FROM   Sailor S, Reserves R  
WHERE  S.sid = R.sid
```

- Naïve, straightforward approach: **combine all tuples**, pick good ones

```
foreach tuple r in R do  
    foreach tuple s in S do  
        if  $r_i == s_j$  then add  $\langle r, s \rangle$  to result
```

- Assume there is no index, R small, S big: better R inner or S?
- What if hash index on S?
- ...this is what cost-based optimization considers!*

Physical Plan Generation

Parser – Checker - Views - Logical plan – Optim1 - **Physical plan** – Optim2 - Execution

- ManyManyMany possible physical query plans for a given logical plan
- physical plan generator tries to select "optimal" one
 - Optimal wrt. response time, throughput
- How are intermediate results passed from children to parents?
 - Temporary files
 - *Evaluate tree bottom-up*
 - *Children write intermediate results to temporary files*
 - *Parents read temporary files*
 - Iterator interface (next)

Sample Query Plan

Parser – Checker - Views - Logical plan – Optim1 - **Physical plan** – Optim2 - Execution

```
SET EXPLAIN ON AVOID_EXECUTE;  
SELECT  C.customer_num, O.order_num  
FROM    customer C, orders O, items I  
WHERE   C.customer_num = O.customer_num  
        AND O.order_num = I.order_num
```

```
for each row in the customer table do:  
  read the row into C  
  for each row in the orders table do:  
    read the row into O  
    if O.customer_num = C.customer_num then  
      for each row in the items table do:  
        read the row into I  
        if I.order_num = O.order_num then  
          accept the row and send to user  
        end if  
      end for  
    end if  
  end for  
end for
```

IBM Informix Dynamic Server

Ex: Iterator for Table Scan

- **open()**
 - Allocate buffer space
- **getNext()**
 - If no block of R has been read yet:
 read first block from disk
 return (R==empty ? null : first tuple in block)
 - If no more tuple left in current block:
 read next block of R from disk
 return (R exhausted ? null : first tuple in block)
 - Return next tuple in block
- **close()**
 - Deallocate buffer space

Ex: Iterator for Nested-Loop Join

Parser – Checker - Views - Logical plan – Optim1 - **Physical plan** – Optim2 - Execution

- **open()**
 - R.open(); S.open();
 - r = R.getNext();

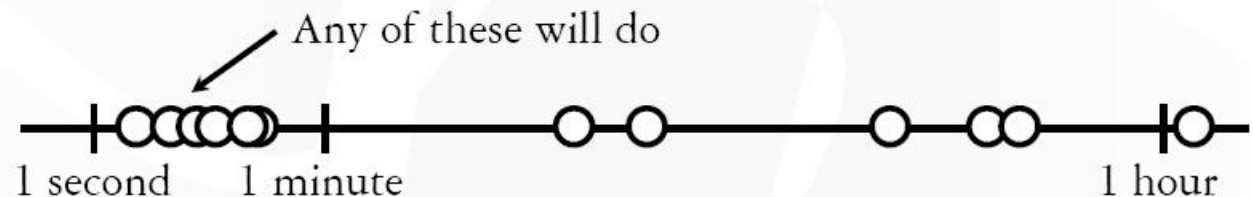
- **getNext()**
 - repeat until r and s join:
 - s = S.getNext();
 - if (s == null)
 - { S.close(); S.open(); s = S.getNext();
 - if (s == null) return null;
 - r = R.getNext();
 - if (r == null) return null;
 - }
 - return <r,s>;

- **close()**
 - R.close(); S.close();

Query Optimization

Parser – Checker - Views - Logical plan – Optim1 - Physical plan – Optim2 - Execution

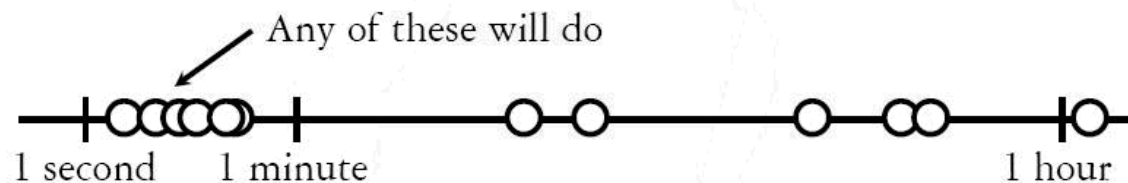
- **Optimization** = find better, equivalent plan
 - Equivalent = produces same result
 - Logical level optimization = aka **heuristic optimization**
 - Physical level optimization = aka **cost-based optimization**
- Two main issues:
 - For a given query, how to **find cheapest plans**?
 - How is **cost** of a plan **estimated**?



(II) Cost-Based Optimization

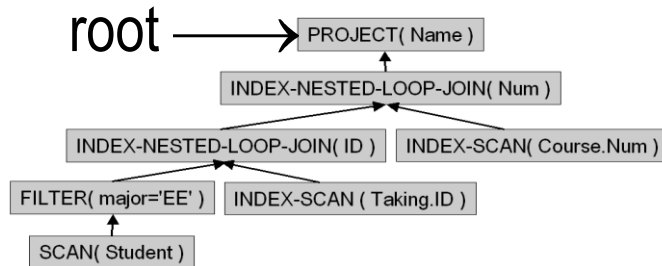
Parser – Checker - Views - Logical plan – Optim1 - Physical plan – **Optim2** - Execution

- logical plan → (efficient) physical plan
- “Cost based” = driven by (estimated) costs of physical situation
 - concrete table sizes, indexes, data distribution, ...
- Approach:
 - **enumerate** all (?) possible physical plans that can be derived from given logical plan
 - **estimate cost** for each plan
 - **pick** best (i.e., least cost) alternative
- **Ideally:** Want to find best plan; **practically:** Avoid worst plans!



Finale: Execution of Tree

Parser – Checker - Views - Logical plan - Rewriter - Physical plan - Optim. - **Execution**



```
result = {};  
root.open();  
do  
{  
    tmp = root.getNext();  
    result += tmp;  
} while (tmp != NULL);  
root.close();  
return result;
```

- Recursive evaluation of tree
 - Requests go down
 - Intermediate result tuples go up
- Often instead: compile into "database machine code" program