

Transaction Management Overview

Ramakrishnan & Gehrke, Chapter 14+

- **Concurrent execution** of user requests is essential for good DBMS performance
 - User requests arrive concurrently
 - Because disk accesses are frequent, and relatively slow, it is important to keep the cpu humming by working on several user programs concurrently
- user's program may carry out many operations on data retrieved, but DBMS only concerned about data read/written from/to database
- A **transaction** (TA) is the DBMS's abstract view of a user program: a sequence of (SQL) reads and writes that is executed as a unit

Concurrency in a DBMS

- Users submit TAs, and can think of each transaction as executing by itself
 - Concurrency achieved by DBMS, which **interleaves** actions (reads/writes of DB objects) of various TAs
 - Each TA must leave the database in a **consistent** state if the DB is consistent when TA begins
 - *DBMS will enforce some ICs, depending on the ICs declared in CREATE TABLE statements.*
 - *Beyond this, the DBMS does not really understand the semantics of the data*
 - *Ex: does not understand how the interest on a bank account is computed*
- Issues:
 - Effect of **interleaving** TAs
 - **crashes**

Atomicity of Transactions

- Two possible TA endings:
 - **commit** after completing all its actions – data must be safe in DB
 - **abort** (by application or DBMS) – must restore original state
- Important property guaranteed by the DBMS: TAs **atomic**
 - Perception: TA executes **all** its actions **in one step**, or **none**
- Technically: DBMS **logs** all actions
 - can **undo** actions of aborted TAs
 - **Write-ahead logging** (WAL): save record of action **before** every update

Transaction Syntax in SQL

- **START TRANSACTION** start TA
- **COMMIT** end TA successfully
- **ROLLBACK** abort TA (undo any changes)
- If none of these TA management commands is present, each statement starts and ends its own TA
 - including all triggers, constraints,...

Anatomy of Conflicts

- Consider two TAs:

```
T1:   BEGIN  A=A-100,  B=B+100  END
T2:   BEGIN  A=1.06*A,  B=1.06*B  END
```

- Intuitively, first TA transfers \$100 from B's account to A's account
 - second TA credits both accounts with a 6% interest payment
- no guarantee that T1 will execute before T2 or vice-versa, if both are submitted together
 - However, net effect must be equivalent to these two TAs running **serially** in some order

Scheduling Transactions: Definitions

- **Serial schedule:**
Schedule that does not interleave the actions of different TAs
- **Equivalent schedules:**
For any database state, the effect (on the set of objects in the database) of executing the first schedule is **identical** to the effect of executing the second schedule
- **Serializable schedule:**
A schedule equivalent to some serial execution of the TAs
- each TA preserves consistency
⇒ every **serializable schedule preserves consistency**




Lock-Based Concurrency Control

- Core issues: What lock modes? What lock conflict handling policy?
- Common lock modes: SX
 - Each TA must obtain an **S** (shared) lock before reading, and an **X** (exclusive) lock before writing
- Lock conflict handling
 - Abort conflicting TA / let it wait / work on previous version
- Locking protocols
 - two-phase locking (strict, non-strict, conservative, ...) – *next!*
 - Timestamp based
 - Multi-version based
 - Optimistic concurrency control

		S	X
S	+	+	-
X	-	-	-

Two-Phase Locking Protocol

- **2PL**
 - All locks acquired before first release (=all locks released after last acquiring)
 - cannot acquire locks after releasing first lock
- allows **only serializable schedules** 😊
 - but complex abort processing
- **Strict 2PL**
 - All locks released when TA completes
- **Strict 2PL simplifies TA aborts** 😊😊



- **Isolation level directives:** summary about TA's intentions, placed **before** TA
 - **SET TRANSACTION READ ONLY**
TA will not write → can be interleaved with other read-only TAs
 - **SET TRANSACTION READ WRITE**
(default)
- assists DBMS optimizer
- Example: Choosing seats in airplane
 - *Find available seat, **reserve** by setting **occ** to **TRUE**; if there is none, abort*
 - *Ask customer for approval. If so, commit, otherwise release seat by setting **occ** to **FALSE**, goto 1*
 - two "TA"s concurrently: can have dirty reads for occ – uncritical! (why?)

Effects of New Isolation Levels

- Consider seat choosing algorithm:
- If run at level **READ COMMITTED**
 - seat choice function will not see seats as booked if reserved but not committed (roll back if over-booked)
 - Repeated queries may yield different seats (other TAs booking in parallel)
- If run at **REPEATABLE READ**
 - any seat found in step 1 will remain available in subsequent queries
 - new tuples entering relation (e.g. switching flight to larger plane) seen by new queries

- **Concurrency control & recovery:** core DBMS functions
- Users need not worry about concurrency
 - System automatically inserts lock/unlocking, schedules TAs, ensures serializability (or what's requested)
- ACID properties!
- Mechanisms:
 - TA scheduling; Strict 2PL !
 - Locks
 - Write-ahead logging (WAL)

Outlook: ACID vs BASE

- **BASE** (Basically Available Soft-state Eventual Consistency)
 - Prefers availability over consistency
 - Relaxing ACID
- **CAP Theorem** [proposed: Eric Brewer; proven: Gilbert & Lynch]:
In a distributed system you can satisfy at most 2 out of the 3 guarantees
 - *Consistency*: all nodes have same data at any time
 - *Availability*: system allows operations all the time
 - *Partition-tolerance*: system continues to work in spite of network partitions
- **Comparison:**
 - Traditional RDBMSs: Strong consistency over availability under a partition
 - Cassandra: Eventual (weak) consistency, availability, partition-tolerance