



JACOBS
UNIVERSITY

NoSQL & NewSQL

Instructors: Peter Baumann

email: p.baumann@jacobs-university.de


tel: -3178

office: room 88, Research 1

With material by Willem Visser

We Don't Want No SQL !

- NoSQL movement: SQL considered slow → only access by id („lookup“)
 - Deliberately abandoning relational world: „too complex“, „not scalable“
 - No clear definition, wide range of systems
 - Values considered black boxes (documents, images, ...)
 - simple operations (ex: key/value storage), horizontal scalability for those
 - ACID → CAP, „eventual consistency“

- Systems
 - Open source  MongoDB, CouchDB, Cassandra, HBase, Riak, Redis
 - Proprietary: Amazon, Oracle, Google , Oracle NoSQL

- See also: <http://glennas.wordpress.com/2011/03/11/introduction-to-nosql-john-nunemaker-presentation-from-june-2010/>

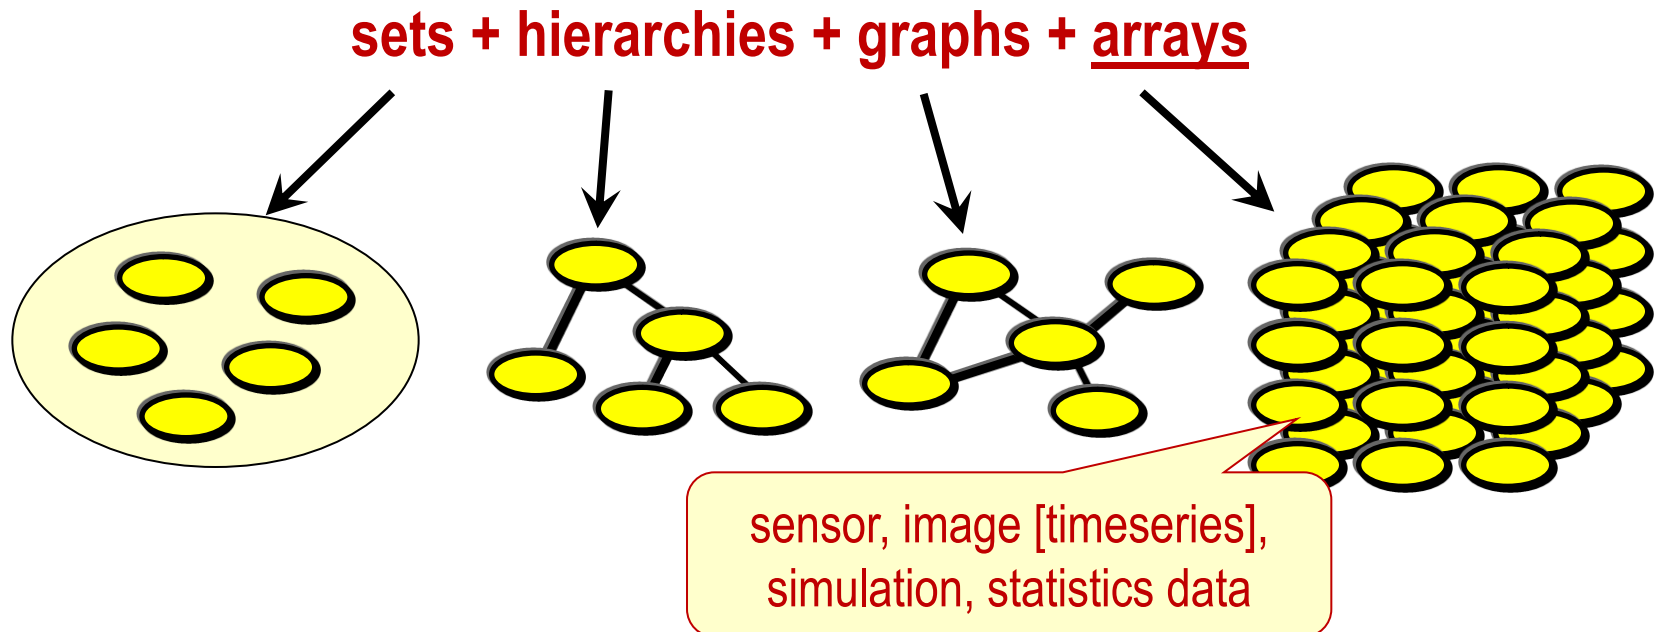
Structural Variety in Big Data

- Stock trading: 1-D sequences (i.e., **arrays**)
- Social networks: large, homogeneous **graphs**
- Ontologies: small, heterogeneous **graphs**
- Climate modelling: 4D/5D **arrays**
- Satellite imagery: 2D/3D **arrays** (+irregularity)
- Genome: long string **arrays**
- Particle physics: **sets** of events
- Bio taxonomies: **hierarchies** (such as XML)
- Documents: key/value stores = **sets** of unique identifiers + whatever
- etc.

Structural Variety in Big Data

- Stock trading: 1-D sequences (i.e., **arrays**)
- Social networks: large, homogeneous **graphs**
- Ontologies: small, heterogeneous **graphs**
- Climate modelling: 4D/5D **arrays**
- Satellite imagery: 2D/3D **arrays** (+irregularity)
- Genome: long string **arrays**
- Particle physics: **sets** of events
- Bio taxonomies: **hierarchies** (such as XML)
- Documents: key/value stores = **sets** of unique identifiers + whatever
- etc.

Structural Variety in Big Data




Ex: Key/Value Store

- Conceptual model: key/value store = **set of key+value**
 - Operations: *Put(key,value)*, *value = Get(key)*
 - → large, distributed **hash table**
- Needed for:
 - twitter.com: tweet id -> information about tweet
 - kayak.com: Flight number -> information about flight, e.g., availability
 - amazon.com: item number -> information about it

- Ex: Cassandra (Facebook; open source)

- Myriads of users, like:




- Like key/value, but value is a **complex document**
- Added: **Search** functionality within document
 - Fulltext search: Lucene/Solr, ElasticSearch...
 - Can support this in architecture, eg, full-text index
- Need: content oriented applications
 - Facebook, Amazon, ...
- Ex: MongoDB, CouchDB

Performance Comparison

- On > 50 GB data:
- MySQL
 - Writes 300 ms avg
 - Reads 350 ms avg
- Cassandra
 - Writes 0.12 ms avg
 - Reads 15 ms avg

How To Make It Fast: 2 x 2 x 2

- 2 kinds of scalability:
 - horizontal scaling over multiple servers
 - vertical scaling for performance on a single server
- Key features needed to achieve this:



- | | |
|---|---|
| ■ For horizontal scaling: | ■ For vertical scaling: |
| <ul style="list-style-type: none">• partition and replicate• automatic failure recovery, database evolution w/o downtime | <ul style="list-style-type: none">• RAM, avoid random disk I/O• minimize overhead for locking & latching, minimize network calls between servers |

[Rick Cattell, <http://www.cattell.net/datastores/ScalabilityRequirements.html>]

Giving Up ACID

- RDBMS provide ACID
- Cassandra provides **BASE**
 - Basically Available Soft-state Eventual Consistency
 - Prefers availability over consistency

CAP Theorem


- Proposed by Eric Brewer, UCB; subsequently proved by Gilbert & Lynch
- In a distributed system you can satisfy at most 2 out of the 3 guarantees
 - **Consistency**: all nodes have same data at any time
 - **Availability**: system allows operations all the time
 - **Partition-tolerance**: system continues to work in spite of network partitions
- Traditional RDBMSs
 - Strong consistency over availability under a partition
- Cassandra
 - Eventual (weak) consistency, Availability, Partition-tolerance

- Previous „young radicals“ approaches subsumed under „NoSQL“
- = we want „no SQL“
- Well..., „not only SQL“
 - After all, a QL is quite handy
 - So, QLs coming into play again (and 2-phase commits = ACID!)
- Ex: MongoDB: „tuple“ = JSON structure

```
db.inventory.find(  
  { type: 'food',  
    $or: [ { qty: { $gt: 100 } }, { price: { $lt: 9.95 } } ]  
  } )
```

Ex 1: Graph Store

- Conceptual model: Labeled, directed, attributed multi-graph
 - Multi-graph = multiple edges between nodes
- Needed by: social networks



[blog.revolutionanalytics.com]


Ex 1: Graph Store



[blog.revolutionanalytics.com]


Ex 1: Graph Store

- Conceptual model: Labeled, directed, attributed multi-graph
 - Multi-graph = multiple edges between nodes
- Needed by: social networks
 - My friends, who has no / many followers, closed communities, new agglomerations, new themes, ...
- Sample system: Neo4j
- Why not relational DB? can model graphs!
 - but “endpoints of an edge” already requires (expensive) join
 - No support for global ops like transitive hull



Ex 2: Array Databases

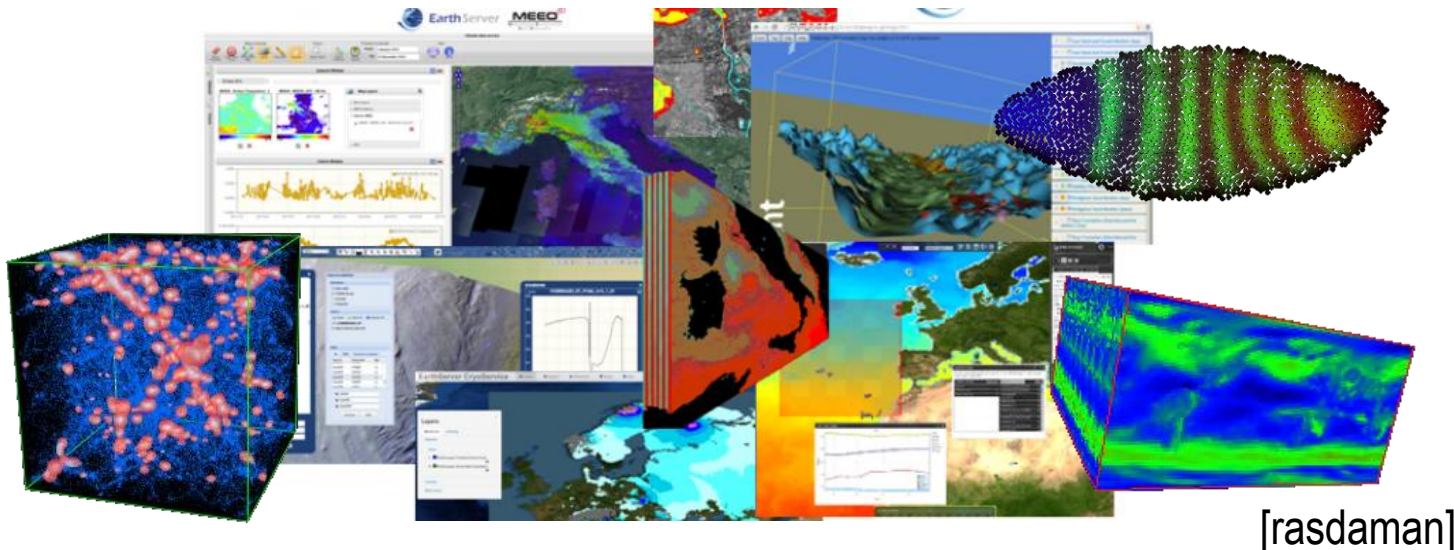

- **Array DBMSs** for declarative queries on massive n-D arrays
 - Ex: **rasdaman** = **Array DBMS** for massive n-D arrays
- ```
select img.green[x0:x1,y0:y1] > 130
from LandsatArchive as img
```
- Array DBMSs can be 200x RDBMS [Cudre-Maroux]
  - Demo at <http://standards.rasdaman.org>





# Array Analytics

- **Array Analytics** :=  
*Efficient analysis on multi-dimensional arrays of a size several orders of magnitude above the evaluation engine's main memory*
- Essential property:  $n$ -D Euclidean neighborhood



# ISO Array SQL

## ■ ISO 9075 Part 15: SQL/MDA

- resolved by ISO SQL WG in June 2014

## ■ n-D arrays as attributes

```
create table LandsatScenes(
 id: integer not null, acquired: date,
 scene: row(band1: integer, ..., band7: integer) array [0:4999,0:4999])
```

## ■ declarative array operations

```
select id, encode(scene.band1-scene.band2)/(scene.nband1+scene.band2)), „image/tiff“)
from LandsatScenes
where acquired between „1990-06-01“ and „1990-06-30“ and
 avg(scene.band3-scene.band4)/(scene.band3+scene.band4)) > 0
```

Information technology — Database languages — SQL —

**Part 15:**

**Multi-Dimensional Arrays (SQL/MDA)**

*Technologies de l'information — Langages de base de données — SQL —*

*Partie 15: Tableaux multi-dimensionnels (SQL/MDA)*

Document type: Technical Report

Document subtype: Technical Report (TR)

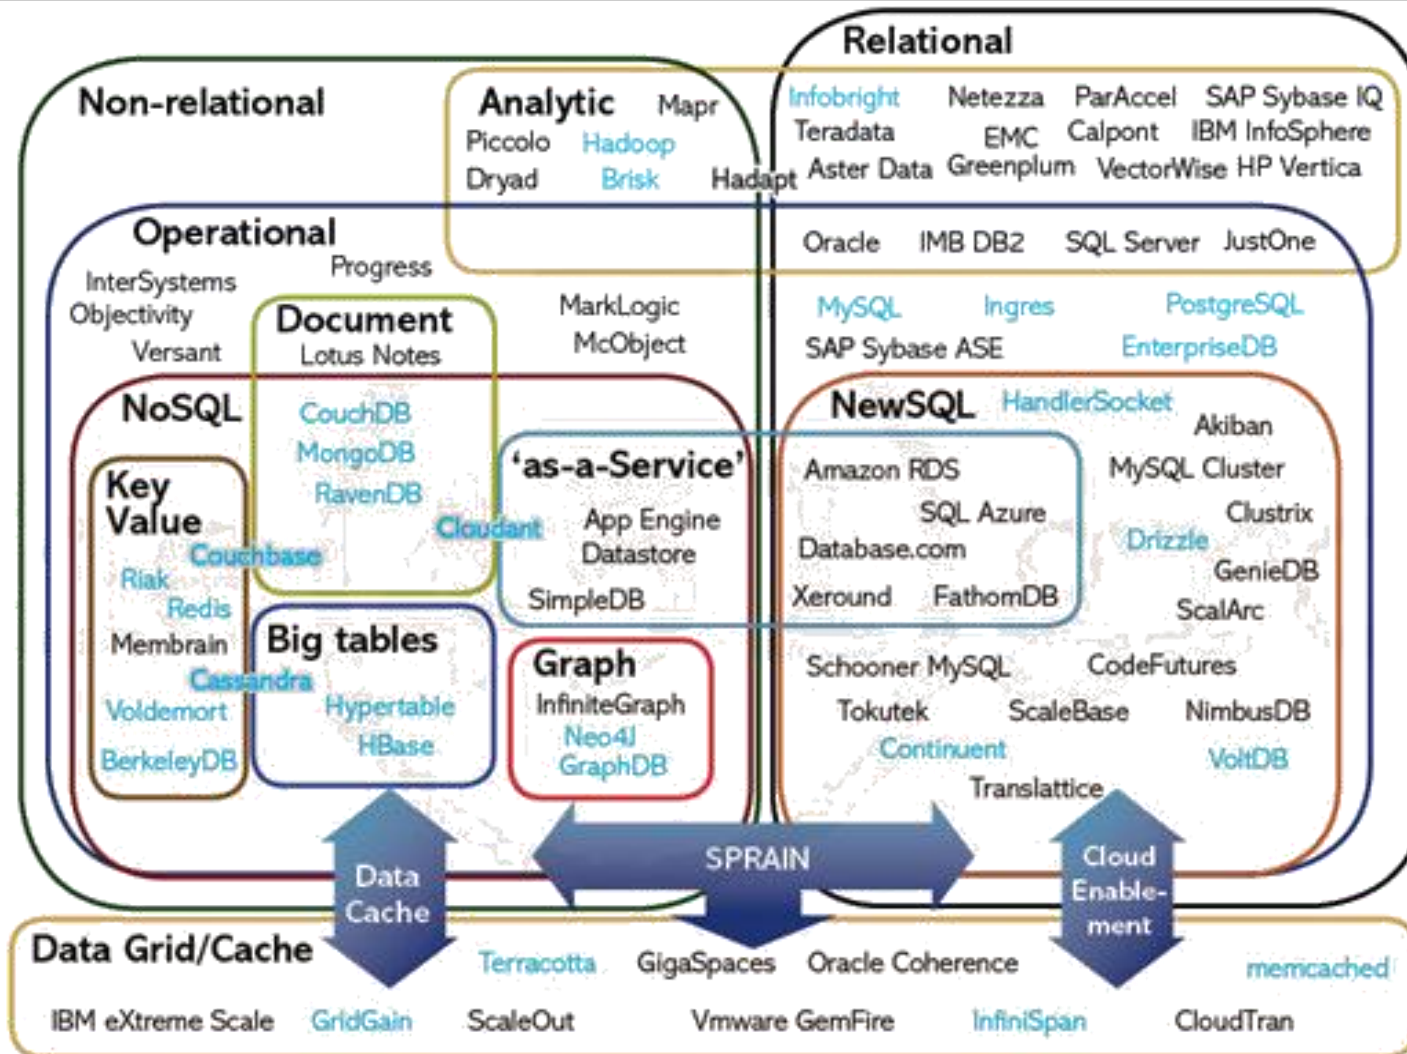
Document stage: (3) CD under Consideration

Document language: English

- Michael Stonebraker: „no one size fits all“
- NoSQL: sacrifice functionality for performance – no QL, only key access
- Swinging back from NoSQL:  
declarative QLs considered good, but SQL often inadequate
- Definition 1: NewSQL = SQL with enhanced performance architectures
- Definition 2: NewSQL = SQL enhanced with, eg, new data types
  - Some call this NoSQL

- *The Relational Empire strikes back*
- Observation: fetching long tuples overhead when few attributes needed
- Brute-force decomposition: one value (plus key)
  - Ex:  $\text{Id+SNLRH} \rightarrow \text{Id+S}, \text{Id+N}, \text{Id+L}, \text{Id+R}, \text{Id+H}$
  - Column-oriented storage: each binary table separate file
- Observation: with clever architecture, reassembly of tuple pays off
- Sample system: MonetDB
  - All major vendors say they have one, but caveat

# The Explosion of DBMSs



[451 group]

...not  
entirely correct



# Database Landscape Map – December 2012

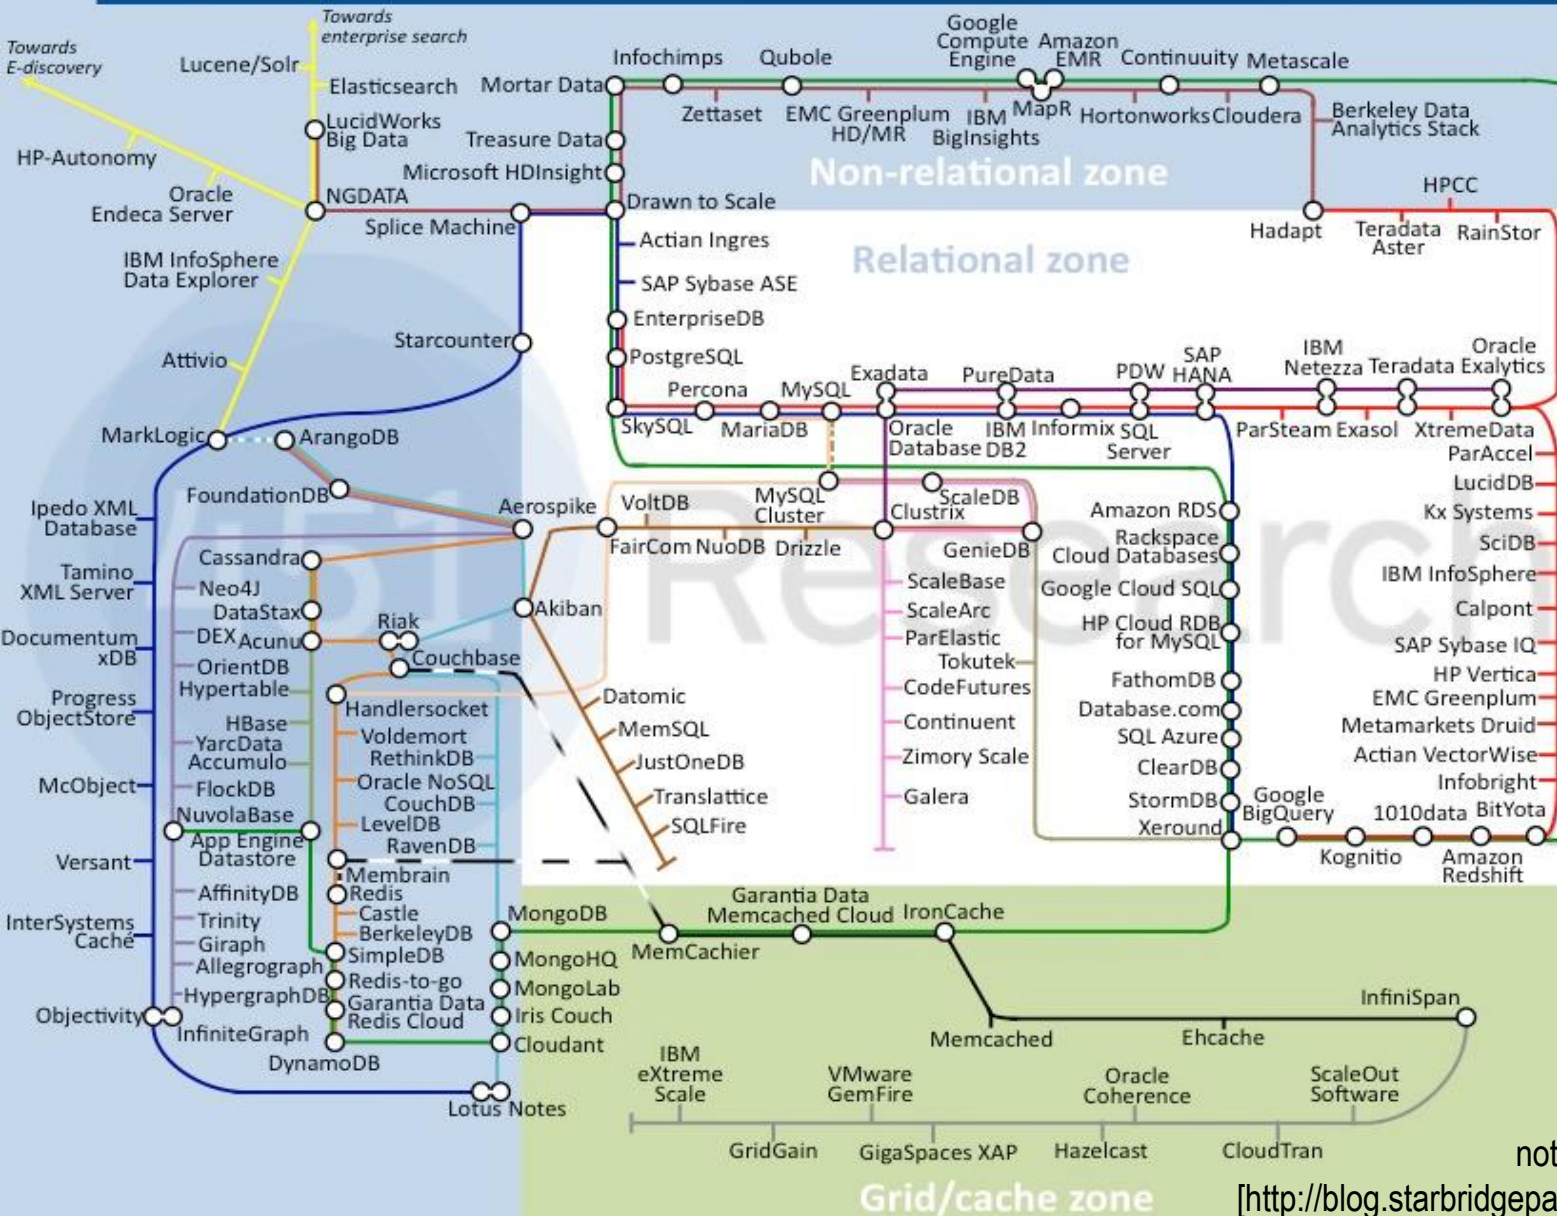
451 Research

## Key:

- Operational
- Analytic
- as-a-Service
- - - NoSQL extension
- BigTables
- Graph
- Document
- Key value stores
- Key value direct access
- Hadoop
- - - NewSQL extension
- Storage engines
- Advanced clustering/sharding
- New SQL databases
- - - Data caching extension
- Data caching
- Data grid
- Index-based data management
- Appliances

www.451research.com

@maslett



not entirely correct/complete

[<http://blog.starbridgepartners.com>, 2013-aug19]

- Fresh approach to scalable data services: NoSQL, NewSQL
  - Diversity of technology → pick best of breed for specific problem
- Avenue 1: **Modular data frameworks** to coexist
  - Heterogeneous model coupling barely understood - needs research
- Avenue 2: concepts **assimilated by relational vendors**
  - Like fulltext, object-oriented, SPARQL, ... cf „Oracle NoSQL“
- “SQL-as-a-service”
  - Amazon RDS, Microsoft SQL Azure, Google Cloud SQL
- *More than ever, experts in data management needed !*
  - *More generally: data engineers*