

# MapReduce

# Why MapReduce?

- Motivation: **Large Scale Data Processing**
- Want to process lots of data ( > 1 TB)
- Want to parallelize across hundreds/thousands of CPUs
- ... Want to make this easy
  - MPI has programming overhead
- MapReduce Idea: simple, highly scalable, generic **parallelization model**
- Automatic parallelization & distribution
- Fault-tolerant
- Clean abstraction for programmers
- status & monitoring tools

# Who Uses MapReduce?

- At Google:
  - Index construction for Google Search
  - Article clustering for Google News
  - Statistical machine translation
- At Yahoo!:
  - “Web map” powering Yahoo! Search
  - Spam detection for Yahoo! Mail
- At Facebook:
  - Data mining
  - Ad optimization

# Overview

- **MapReduce**: the concept
- **Hadoop**: the implementation
- Query Languages for Hadoop
- **Spark**: the improvement
- MapReduce vs databases
- Conclusion

# MapReduce: the concept

## Credits:

- David Maier
- Google
- Shiva Teja Reddi Gopidi

# Preamble:

## Merits of Functional Programming (FP)

- FP: input determines output – *and nothing else*
  - No other knowledge used (global variables!)
  - No other data modified (global variables!)
  - Every function invocation generates new data
- Opposite: procedural programming → side effects
  - Unforeseeable interference between parallel processes  
→ difficult/impossible to ensure deterministic result
- (function, value set) must be monoid
- Advantage of FP: parallelization can be arranged automatically
  - can (automatically!) reorder or parallelize execution - data flow implicit

being side-effect free is  
a key property of SQL

# Programming Model

- Goals: large data sets, processing distributed over 1,000s of nodes
  - Abstraction to express simple computations
  - Hide details of parallelization, data distribution, fault tolerance, load balancing
    - MapReduce engine performs all housekeeping
- Inspired by primitives from functional PLs like Lisp, Scheme, Haskell
- Input, output are sets of key/value pairs
- Users implement interface of two functions:

`map (inKey, inValue) -> (outKey, intermediateValuelist )`

← aka „group by“ in SQL


`reduce(outKey, intermediateValuelist) -> outValuelist`

← aka aggregation in SQL

# Ex 1: Count Word Occurrences

```
map(String inKey, String inValue):
    // inKey: document name
    // inValue: document contents
    for each word w in inValue:
        EmitIntermediate(w, "1");
```

```
reduce(String outputKey, Iterator auxValues):
    // outKey: a word
    // outValues: a list of counts
    int result = 0;
    for each v in auxValues:
        result += ParseInt(v);
    Emit( AsString(result) );
```




[image: Google]



# Ex 2: Distributed Grep

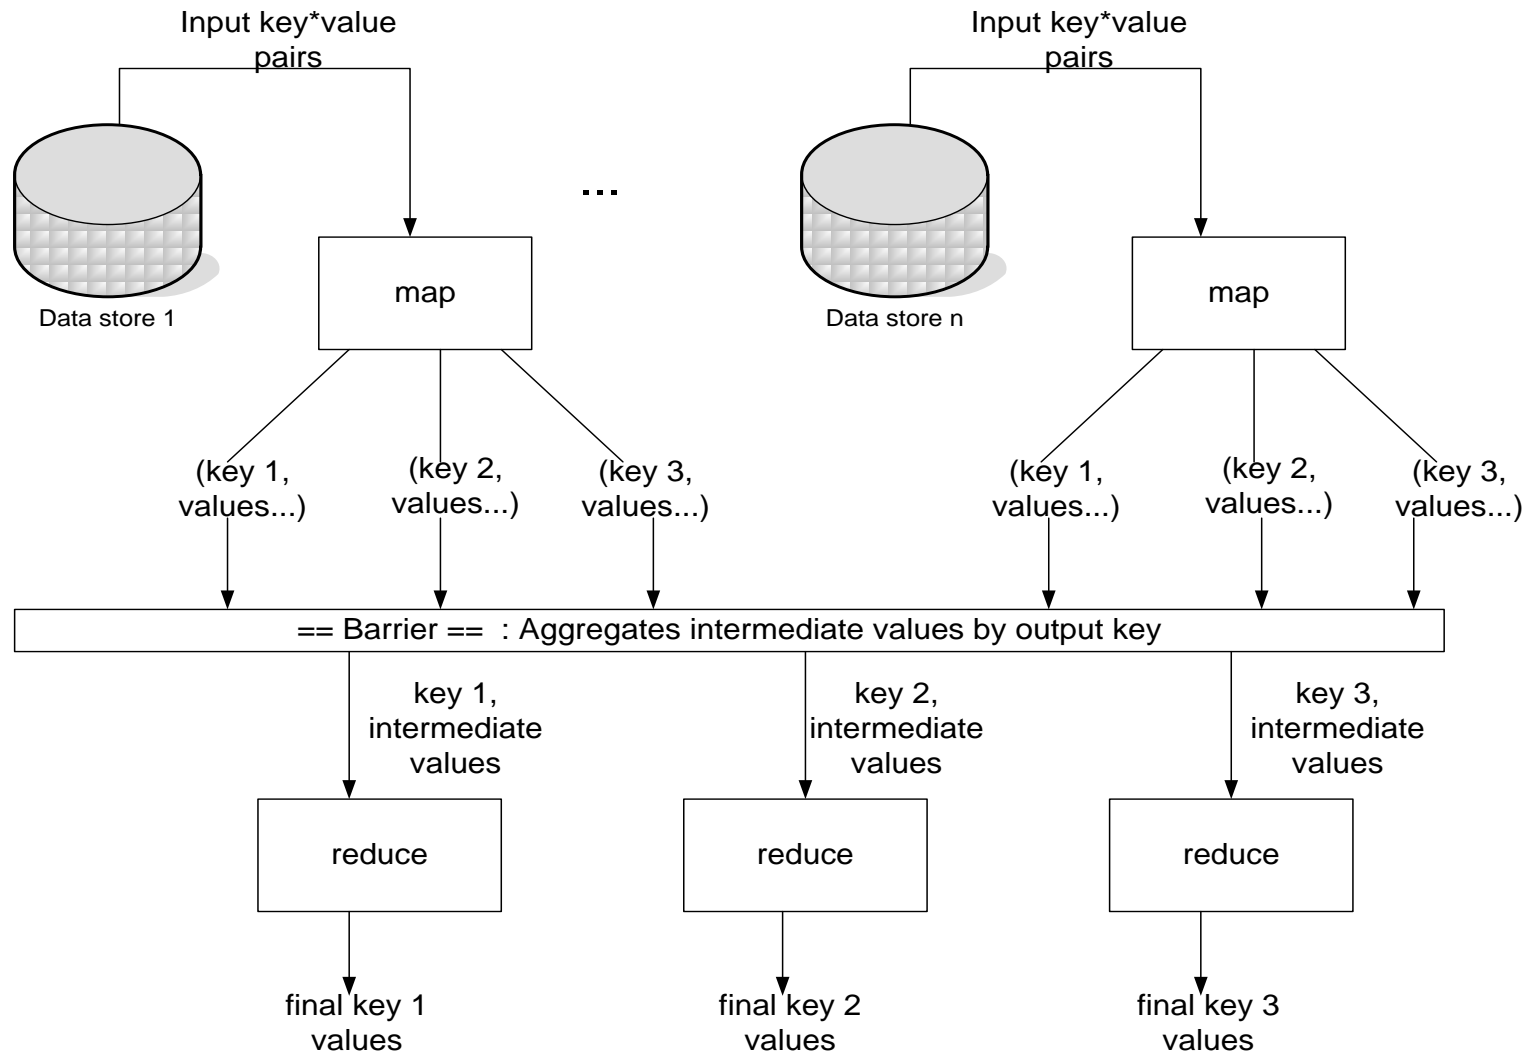
- map function **emits line** if matches given pattern
  - identity function that just copies supplied intermediate data to output
- Application 1: Count of **URL Access Frequency**
  - logs of web page requests → **map()** → <URL,1>
  - all values for same URL → **reduce()** → <URL, total count>
- Application 2: Inverted **Index**
  - Document → **map()** → sequence of <word, document ID> pairs
  - all pairs for a given word → **reduce()** sorts document IDs → <word, list(document ID)>
  - set of all output pairs = simple inverted index
  - easy to extend for word positions

# Ex 3: Relational Join



- Map function M: “hash on key attribute”:  $(?, \text{tuple}) \rightarrow \text{list}(\text{key}, \text{tuple})$
- Reduce function R: “join on each k value”:  $(\text{key}, \text{list}(\text{tuple})) \rightarrow \text{list}(\text{tuple})$

# Map & Reduce



# Map Reduce Patent

- Google granted US Patent 7,650,331, January 2010
- **System and method for efficient large-scale data processing**  
A large-scale data processing system and method includes one or more application-independent map modules configured to read input data and to apply at least one **application-specific map operation** to the input data to produce intermediate data values, wherein the map operation is automatically parallelized across multiple processors in the parallel processing environment. A plurality of intermediate data structures are used to store the intermediate data values. One or more application-independent reduce modules are configured to retrieve the intermediate data values and to apply at least one **application-specific reduce operation** to the intermediate data values to provide output data.



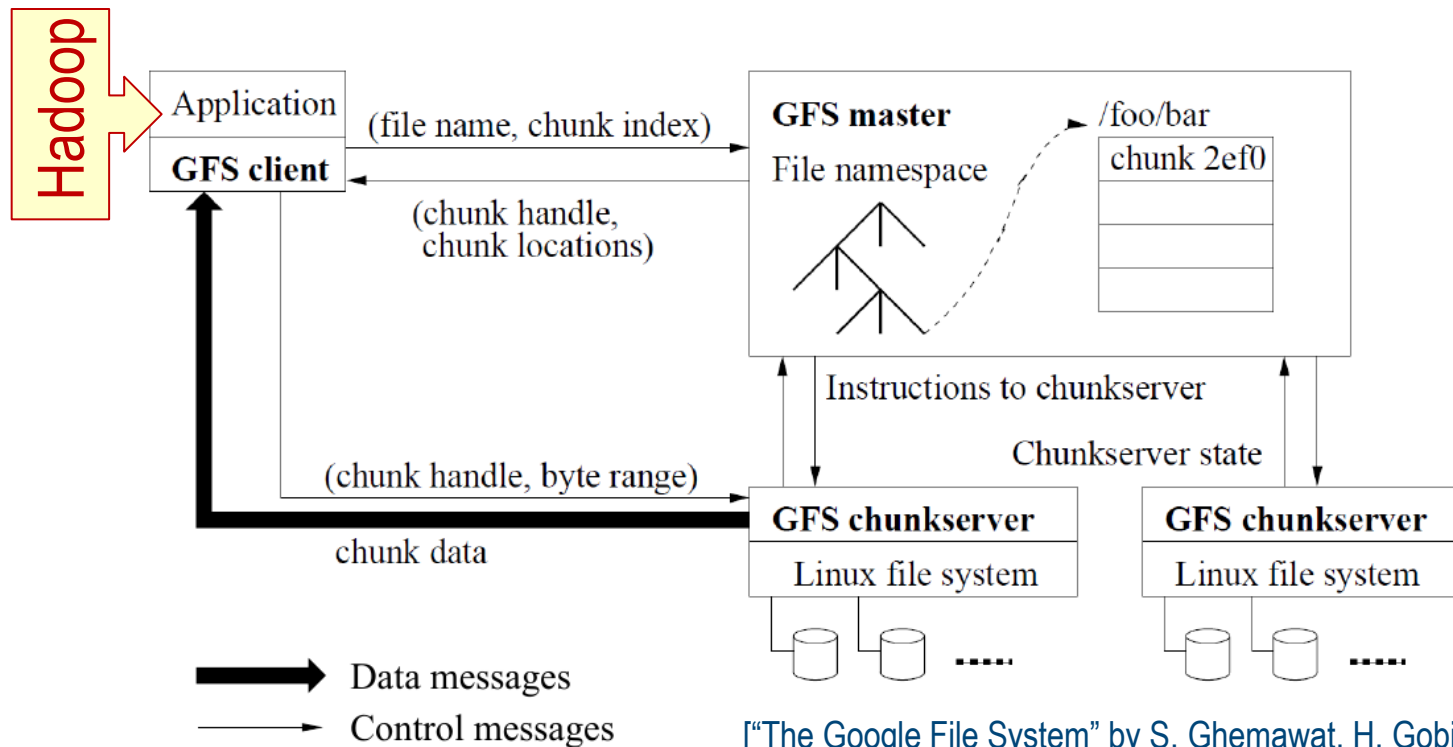
# Hadoop: a MapReduce implementation

## Credits:

- David Maier, U Wash
- Costin Raiciu
- “The Google File System” by S. Ghemawat, H. Gobioff, and S.-T. Leung, 2003
- [https://hadoop.apache.org/docs/r1.0.4/hdfs\\_design.html](https://hadoop.apache.org/docs/r1.0.4/hdfs_design.html)

# Hadoop Distributed File System

- HDFS = scalable, fault-tolerant file system
  - modeled after Google File System (GFS)
  - 64 MB blocks („chunks“)



["The Google File System" by S. Ghemawat, H. Gobioff, and S.-T. Leung, 2003]

# GFS

- Goals:
  - Many inexpensive commodity components – failures happen routinely
  - Optimized for small # of large files (ex: a few million of 100+ MB files)
- relies on **local storage** on each node
  - parallel file systems: typically dedicated I/O servers (ex: IBM GPFS)
- metadata (file-chunk mapping, replica locations, ...) in **master node's** RAM
  - Operation log on master's local disk, replicated to remotes → master crash recovery!
  - „Shadow masters“ for read-only access

## HDFS differences?

- No random write; append only
- Implemented in Java, emphasizes platform independence
- terminology: namenode → master, block → chunk, ...

# GFS Consistency

- Relaxed consistency model
  - tailored to Google's highly distributed applications, simple & efficient to implement
- File **namespace** mutations are atomic
  - handled exclusively by master; locking guarantees atomicity & correctness
  - master's log defines global total order of operations
- State of **file region** after data mutation
  - **consistent**: all clients always see same data, regardless of replica they read from
  - **defined**: consistent, plus all clients see the entire data mutation
  - **undefined but consistent**: result of concurrent successful mutations; all clients see same data, but may not reflect any one mutation
  - **inconsistent**: result of a failed mutation




# GFS Consistency: Consequences

- Implications for applications
  - *better not distribute records across chunks!*
  - rely on appends rather than overwrites
  - **application-level** checksums, checkpointing, writing self-validating & self-identifying records
- Typical use cases (or “hacking around relaxed consistency”)
  - writer generates file from beginning to end and then atomically renames it to a permanent name under which it is accessed
  - writer inserts periodical checkpoints, readers only read up to checkpoint
  - many writers concurrently append to file to merge results, reader skip occasional padding and repetition using checksums

# Replica Placement

- Goals of placement policy
  - scalability, reliability and availability, maximize network bandwidth utilization
- Background: GFS clusters are highly distributed
  - 100s of chunkservers across many racks
  - accessed from 100s of clients from same or different racks
  - traffic between machines on different racks may cross many switches
  - bandwidth between racks typically lower than within rack



# Replica Placement

- Goals of placement policy
  - scalability, reliability and availability, maximize network bandwidth utilization
- Background: GFS clusters are highly distributed
  - 100s of chunkservers across many racks
  - accessed from 100s of clients from same or different racks
  - traffic between machines on different racks may cross many switches
  - bandwidth between racks typically lower than within rack
- Selecting a chunkserver
  - place chunks on servers with below-average disk space utilization
  - place chunks on servers with low number of recent writes
  - spread chunks across racks (see above)

# Hadoop Job Management Framework

- **JobTracker** = daemon service for submitting & tracking MapReduce jobs
- **TaskTracker** = slave node daemon in the cluster accepting tasks (Map, Reduce, & Shuffle operations) from a JobTracker

## Discussion:

- Pro: replication & automated restart of failed tasks  
→ highly reliable & available
- Con: 1 Job Tracker per Hadoop cluster, 1 Task Tracker per slave node  
→ single point of failure

# Optimizations / 1

- Problem:  
No reduce can start until map is complete  
→ single slow disk controller can rate-limit whole process
- Solution:  
Master **redundantly** executes slow (“straggler”) map tasks;  
uses results of first copy to finish
- *Why is it safe to redundantly execute map tasks?  
Wouldn't this mess up the total computation?*

# Optimizations / 2

- Problem:  
excessive data transport between map() and reduce() workers
- Approach:  
“Combiner” functions can run on same machine as a mapper
  - “mini-reduce phase” followed by “final” reduce phase
  - saves bandwidth
- *Under what conditions is it sound to use a combiner?*

# Discussion

- MapReduce concept:
  - One-input two-stage data flow extremely rigid
  - Most suitable for independent data
    - *Good: word count*
    - *Not optimal: join, graphs, arrays, ...*
  - HDFS assumes shared-nothing & locality, but datacenters often run SANs
  - (Well-known) algorithms need cumbersome rewriting = special-skill programming
    - *Query frontends: Pig Latin, Hive, etc.*
  - `map()`, `reduce()` procedural Java code  
→ hard to optimize
- Hadoop implementation:
  - All intermediate data communicated via disk
  - Task scheduler: central point of failure
  - HDFS not standards conformant (eg, POSIX)

# Query Languages for MapReduce

Credits:

- Matei Zaharia



# Motivation

- MapReduce is powerful
  - many algorithms can be expressed as a series of MR jobs
- But fairly low-level
  - must think about keys, values, partitioning, etc.
- Can we capture common “job patterns”?
  - Like eg SQL does

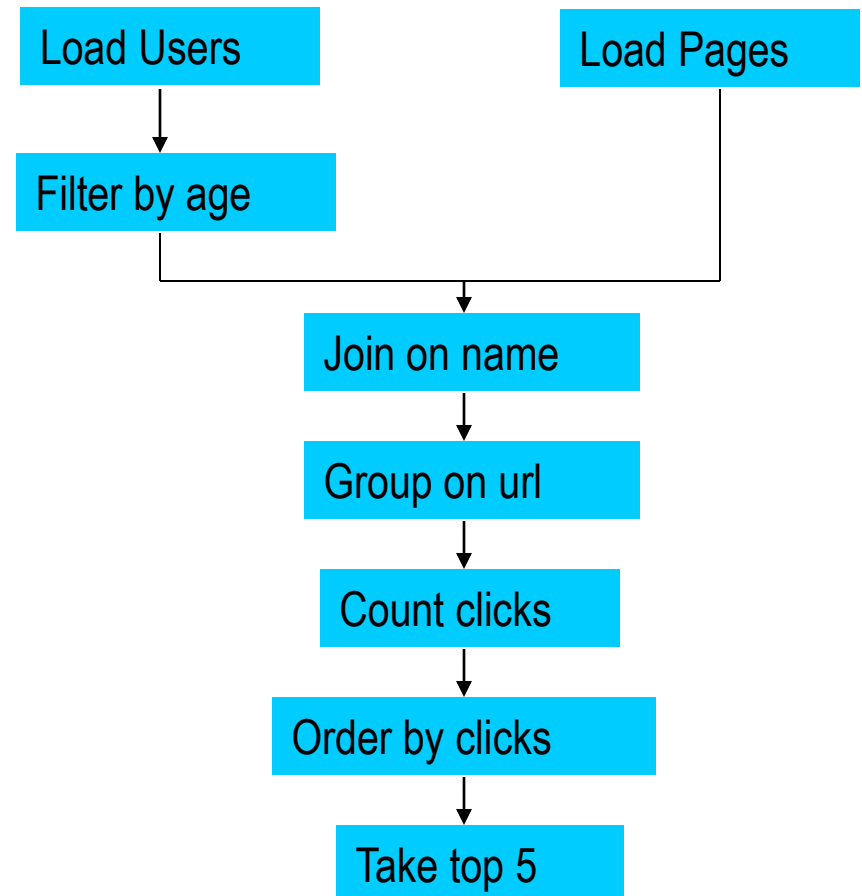
# Pig

- Started at Yahoo! Research
  - Runs about 50% of Yahoo!'s jobs
- Features:
  - Expresses sequences of MapReduce jobs
  - Data model: nested “bags” of items
  - Provides relational (SQL) operators (JOIN, GROUP BY, etc)
  - Easy to plug in Java functions



# Example Problem

- user data in one file
- website data in another
- find top 5 most visited pages
- by users aged 18-25



[<http://wiki.apache.org/pig-data/attachments/PigTalksPapers/attachments/ApacheConEurope09.ppt>]

# In MapReduce

```
import java.io.IOException;
import java.util.ArrayList;
import java.util.Iterator;
import java.util.List;

import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.io.Writable;
import org.apache.hadoop.io.WritableComparable;
import org.apache.hadoop.mapred.FileInputFormat;
import org.apache.hadoop.mapred.FileOutputFormat;
import org.apache.hadoop.mapred.JobConf;
import org.apache.hadoop.mapred.KeyValueTextInputFormat;
import org.apache.hadoop.mapred.Mapper;
import org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.RecordReader;
import org.apache.hadoop.mapred.Reducer;
import org.apache.hadoop.mapred.Reporter;
import org.apache.hadoop.mapred.SequenceFileInputFormat;
import org.apache.hadoop.mapred.SequenceFileOutputFormat;
import org.apache.hadoop.mapred.TextInputFormat;
import org.apache.hadoop.mapred.JobControl;
import org.apache.hadoop.mapred.JobControl.JobControl;
import org.apache.hadoop.mapred.lib.IdentityMapper;

public class MRExample {
    public static class LoadPages extends MapReduceBase
        implements Mapper<LongWritable, Text, Text> {

        public void map(LongWritable k, Text val,
            OutputCollector<Text, Text> oc,
            Reporter reporter) throws IOException {
            // Pull the key out
            String line = val.toString();
            int firstComma = line.indexOf(',');
            String key = line.substring(0, firstComma);
            String value = line.substring(firstComma + 1);
            Text outKey = new Text(key);
            // Prepend an index to the value so we know which file
            // it came from.
            Text outVal = new Text("1" + value);
            oc.collect(outKey, outVal);
        }
    }

    public static class LoadAndFilterUsers extends MapReduceBase
        implements Mapper<LongWritable, Text, Text> {

        public void map(LongWritable k, Text val,
            OutputCollector<Text, Text> oc,
            Reporter reporter) throws IOException {
            // Pull the key out
            String line = val.toString();
            int firstComma = line.indexOf(',');
            String value = line.substring(firstComma + 1);
            int age = Integer.parseInt(value);
            if (age < 18 || age > 25) return;
            String key = line.substring(0, firstComma);
            Text outKey = new Text(key);
            // Prepend an index to the value so we know which file
            // it came from.
            Text outVal = new Text("2" + value);
            oc.collect(outKey, outVal);
        }
    }

    public static class Join extends MapReduceBase
        implements Reducer<Text, Text, Text> {

        public void reduce(Text key,
            Iterator<Text> iter,
            OutputCollector<Text, Text> oc,
            Reporter reporter) throws IOException {
            // For each value, figure out which file it's from and
            // accordingly.
            List<String> first = new ArrayList<String>();
            List<String> second = new ArrayList<String>();

            while (iter.hasNext()) {
                Text t = iter.next();
                String value = t.toString();
                if (value.charAt(0) == '1')
                    first.add(value.substring(1));
                else second.add(value.substring(1));
            }

            reporter.setStatus("OK");
        }
    }

    // Do the cross product and collect the values
    for (String s1 : first) {
        for (String s2 : second) {
            String outval = key + "," + s1 + "," + s2;
            oc.collect(null, new Text(outval));
            reporter.setStatus("OK");
        }
    }
}

public static class LoadJoined extends MapReduceBase
    implements Mapper<Text, Text, Text, LongWritable> {

    public void map(
        Text k,
        Text val,
        OutputCollector<Text, LongWritable> oc,
        Reporter reporter) throws IOException {
        // Find the url
        String line = val.toString();
        int firstComma = line.indexOf(',');
        int secondComma = line.indexOf(',', firstComma);
        String key = line.substring(firstComma, secondComma);
        // drop the rest of the record, I don't need it anymore,
        // just pass a 1 for the combiner/reducer to sum instead.
        Text outKey = new Text(key);
        Text outVal = new Text("1");
        oc.collect(outKey, new LongWritable(1L));
    }
}

public static class ReduceUrls extends MapReduceBase
    implements Reducer<Text, LongWritable, WritableComparable,
        Writable> {

    public void reduce(
        Text key,
        Iterator<LongWritable> iter,
        OutputCollector<WritableComparable, Writable> oc,
        Reporter reporter) throws IOException {
        // Add up all the values we see
        long sum = 0;
        while (iter.hasNext()) {
            sum += iter.next().get();
            reporter.setStatus("OK");
        }
        oc.collect(key, new LongWritable(sum));
    }
}

public static class LoadClicks extends MapReduceBase
    implements Mapper<WritableComparable, Writable, LongWritable,
        Text> {

    public void map(
        WritableComparable key,
        Writable val,
        OutputCollector<LongWritable, Text> oc,
        Reporter reporter) throws IOException {
        oc.collect((LongWritable)val, (Text)key);
    }
}

public static class LimitClicks extends MapReduceBase
    implements Reducer<LongWritable, Text, LongWritable, Text> {

    int count = 0;
    public void reduce(
        LongWritable key,
        Iterator<Text> iter,
        OutputCollector<LongWritable, Text> oc,
        Reporter reporter) throws IOException {
        // Only output the first 100 records
        while (count < 100 && iter.hasNext()) {
            oc.collect(key, iter.next());
            count++;
        }
    }
}

public static void main(String[] args) throws IOException {
    JobConf jp = new JobConf(MRExample.class);
    jp.setJobName("Load Pages");
    jp.setMapperClass(LoadPages.class);
    jp.setInputFormatClass(KeyValueTextInputFormat.class);
    jp.setOutputFormatClass(LongWritable.class);
    jp.setCombinerClass(ReduceUrls.class);
    jp.setReducerClass(ReduceUrls.class);
    jp.setOutputPath(new Path("/user/gates/tmp/indexed_pages"));
    jp.setNumReduceTasks(0);
    Job loadPages = new Job(jp);

    JobConf jfu = new JobConf(MRExample.class);
    jfu.setJobName("Load and Filter Users");
    jfu.setInputFormatClass(KeyValueTextInputFormat.class);
    jfu.setOutputFormatClass(LongWritable.class);
    jfu.setMapperClass(LoadAndFilterUsers.class);
    jfu.setInputFormatClass(KeyValueTextInputFormat.class);
    jfu.setOutputPath(new Path("/user/gates/tmp/filtered_users"));
    jfu.setNumReduceTasks(0);
    Job loadUsers = new Job(jfu);

    JobConf join = new JobConf(MRExample.class);
    join.setJobName("Join Users and Pages");
    join.setInputFormatClass(KeyValueTextInputFormat.class);
    join.setOutputFormatClass(LongWritable.class);
    join.setMapperClass(Join.class);
    join.setCombinerClass(ReduceUrls.class);
    join.setReducerClass(ReduceUrls.class);
    join.setOutputPath(new Path("/user/gates/tmp/joined"));
    join.setNumReduceTasks(50);
    Job joinJob = new Job(join);
    joinJob.addDependingJob(loadPages);
    joinJob.addDependingJob(loadUsers);

    JobConf group = new JobConf(MRExample.class);
    group.setJobName("Group URLs");
    group.setInputFormatClass(KeyValueTextInputFormat.class);
    group.setOutputFormatClass(LongWritable.class);
    group.setMapperClass(LoadClicks.class);
    group.setCombinerClass(ReduceUrls.class);
    group.setReducerClass(ReduceUrls.class);
    group.setOutputPath(new Path("/user/gates/tmp/grouped"));
    group.setNumReduceTasks(50);
    Job groupJob = new Job(group);
    groupJob.addDependingJob(joinJob);

    JobConf top100 = new JobConf(MRExample.class);
    top100.setJobName("Top 100 sites");
    top100.setInputFormatClass(SequenceFileInputFormat.class);
    top100.setOutputFormatClass(LongWritable.class);
    top100.setMapperClass(LimitClicks.class);
    top100.setCombinerClass(LimitClicks.class);
    top100.setReducerClass(LimitClicks.class);
    top100.setOutputPath(new Path("/user/gates/top100sitesforusers18to25"));
    top100.setNumReduceTasks(1);
    Job limit = new Job(top100);
    limit.addDependingJob(groupJob);

    JobControl jc = new JobControl("Find top 100 sites for users
        18 to 25");
    jc.addJob(loadPages);
    jc.addJob(loadUsers);
    jc.addJob(joinJob);
    jc.addJob(groupJob);
    jc.addJob(limit);
    jc.run();
}
```

[<http://wiki.apache.org/pig-data/attachments/PigTalksPapers/attachments/ApacheConEurope09.ppt>]

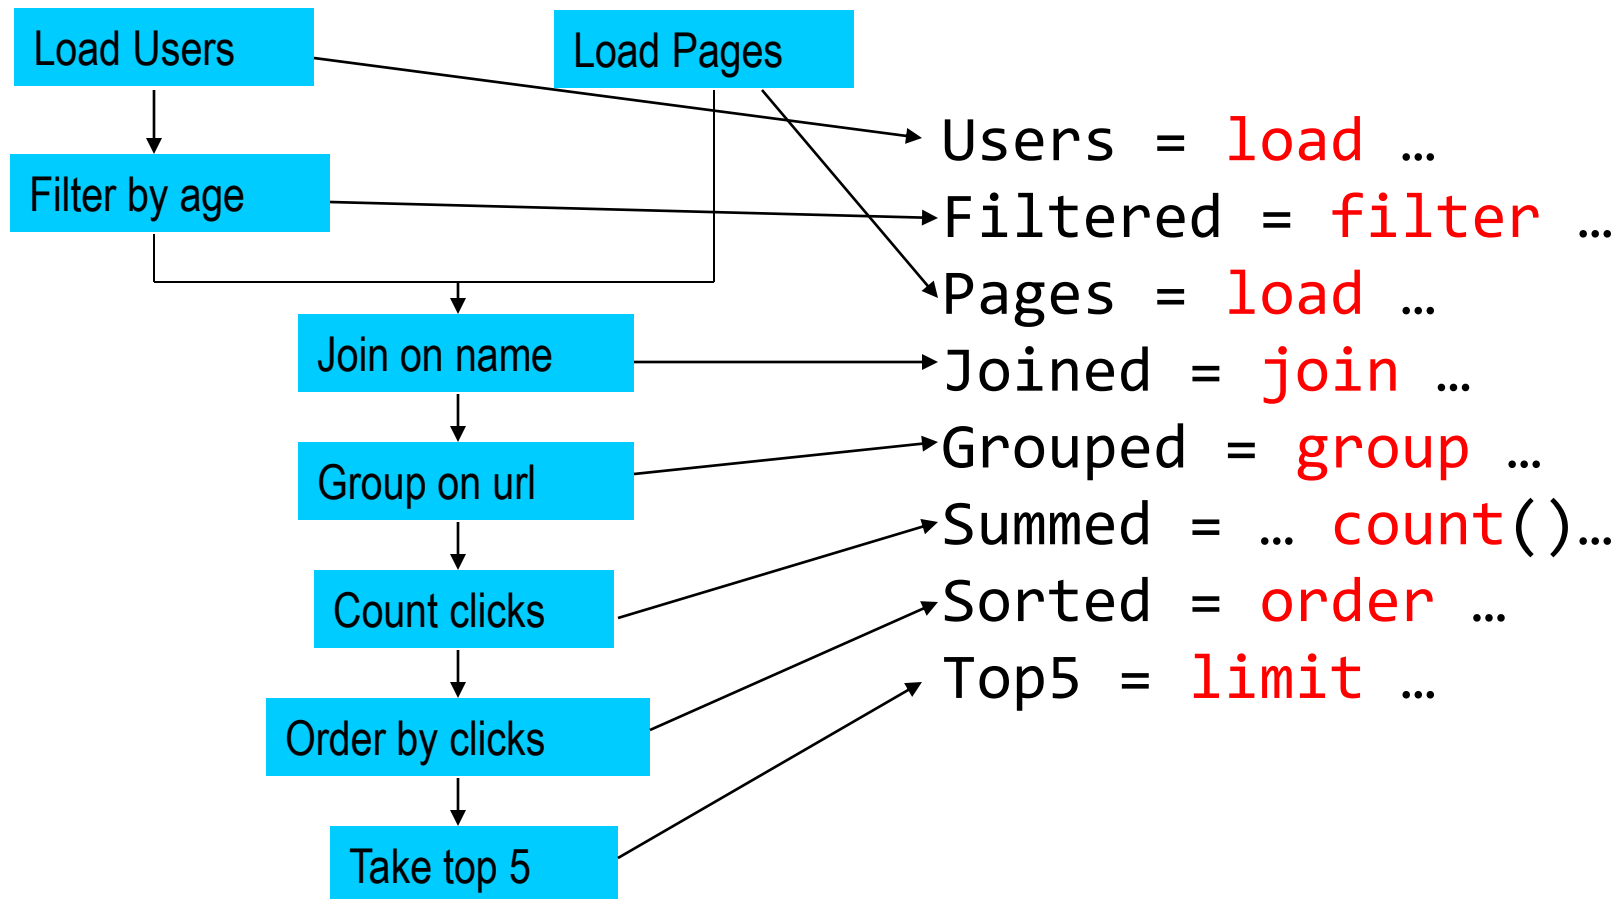
# In Pig Latin

```
Users      = load 'users' as (name, age);
Filtered   = filter Users by
              age >= 18 and age <= 25;
Pages      = load 'pages' as (user, url);
Joined     = join Filtered by name, Pages by user;
Grouped    = group Joined by url;
Summed     = foreach Grouped generate group,
              count(Joined) as clicks;
Sorted     = order Summed by clicks desc;
Top5       = limit Sorted 5;

store Top5 into 'top5sites';
```

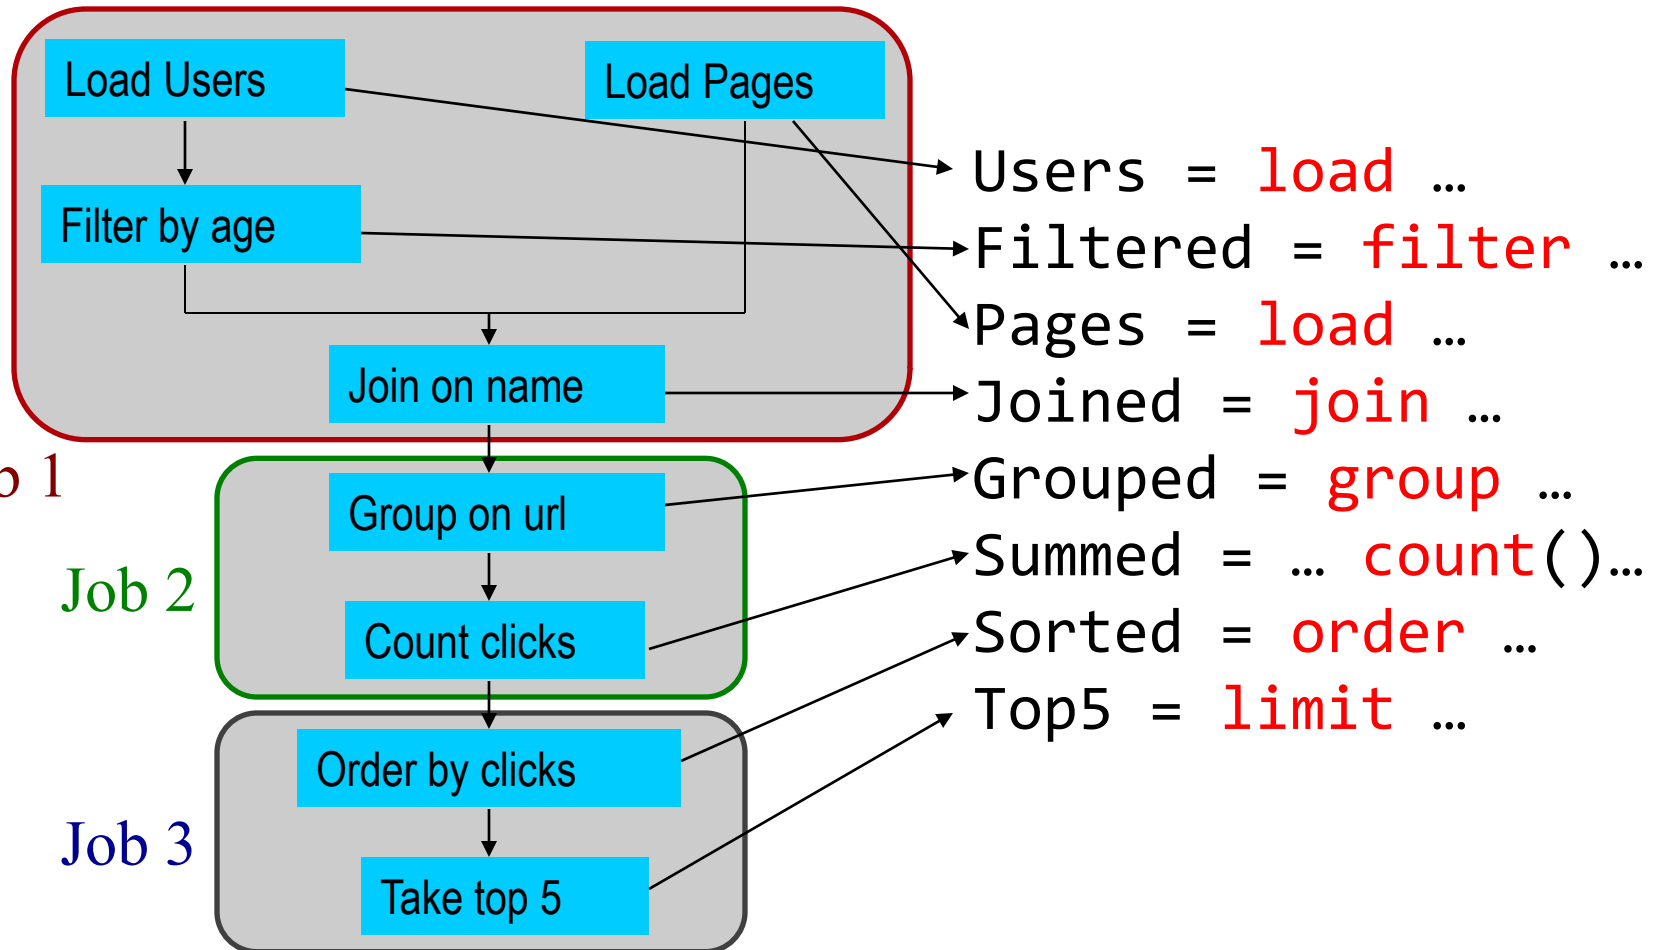
# Translation to MapReduce

Quite natural translation of job components into Pig Latin:



# Translation to MapReduce

Quite natural translation of job components into Pig Latin:



# Hive

- Relational database built on Hadoop

- table schemas, SQL-like query language

```
SELECT word, count(1) AS count
FROM (SELECT explode(split(line, '\s')) AS word
      FROM docs) temp
GROUP BY word
ORDER BY word
```

- can call Hadoop Streaming scripts

- Common relational features:

- table partitioning, complex data types, sampling
- some query optimization

- Developed at Facebook, now Apache

- Today: „data warehouse infrastructure“






# MapReduce vs (Relational) Databases

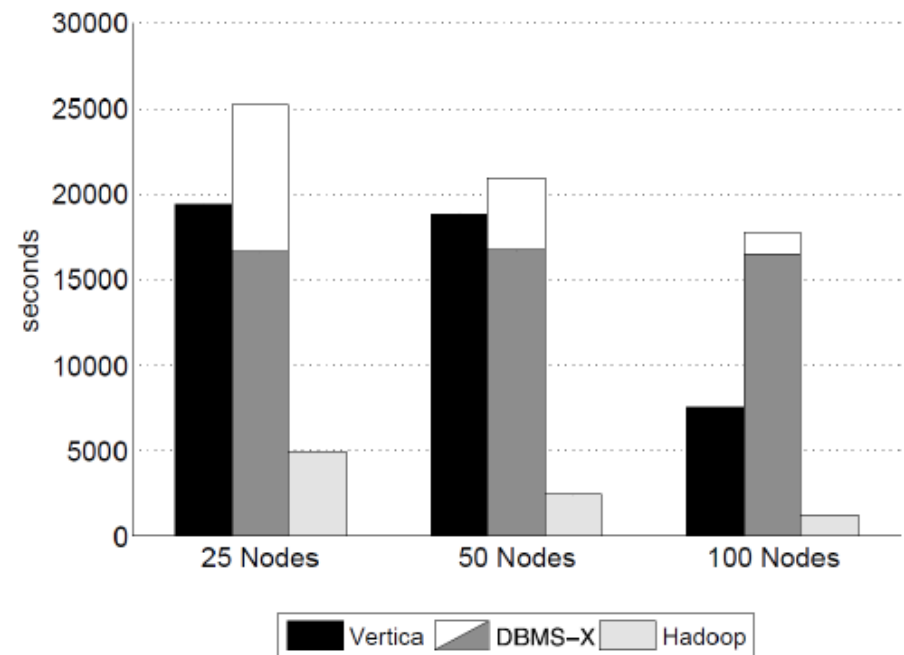
Credits: David Maier

# Grep Task: Load Times

**535 MB/node**




**1 TB/cluster**




[“A Comparison of Approaches to Large-Scale Data Analysis” by A. Pavlo et al., 2004]

# Grep Task: Execution Times

**535 MB/node**



**1 TB/cluster**



[“A Comparison of Approaches to Large-Scale Data Analysis” by A. Pavlo et al., 2004]

# MapReduce Criticism

- Efficiency
  - master makes  $O(M + R)$  scheduling decisions
  - master stores  $O(M * R)$  states in memory
- “Why not use a parallel DBMS instead?”
  - map/reduce is a “giant step backwards”
  - no schema, no indexes, no high-level language
  - not novel at all
  - does not provide features of traditional DBMS
  - incompatible with DBMS tools

# Analytics Tasks

```
CREATE TABLE Documents (  
    url VARCHAR(100)  
    PRIMARY KEY,  
    contents TEXT );
```

```
CREATE TABLE Rankings (  
    pageURL VARCHAR(100)  
    PRIMARY KEY,  
    pageRank INT,  
    avgDuration INT );
```


```
CREATE TABLE UserVisits (  
    sourceIP VARCHAR(16),  
    destURL VARCHAR(100),  
    visitDate DATE,  
    adRevenue FLOAT,  
    userAgent VARCHAR(64),  
    countryCode VARCHAR(3),  
    languageCode  
    VARCHAR(3),  
    searchWord VARCHAR(32),  
    duration INT );
```

## ■ Data set

- 600K unique HTML documents
- 155M user visit records (20 GB/node)
- 18M ranking records (1 GB/node)

[“A Comparison of Approaches to Large-Scale Data Analysis” by A. Pavlo et al., 2004]

# Select Task



## SQL Query:

```
SELECT pageURL, pageRank
FROM Rankings
WHERE pageRank > X
```

## Relational DBMS

- use index on pageRank column
- Relative performance degrades as number of nodes increases

## Hadoop start-up cost increase with cluster size


[“A Comparison of Approaches to Large-Scale Data Analysis” by A. Pavlo et al., 2004]

# Aggregation Task

*“total ad revenue for each source IP, based on user visits table”*


## Variant 1: 2.5M groups

```
SELECT sourceIP, SUM(adRevenue)
FROM UserVisits
GROUP BY sourceIP
```



## Variant 2: 2,000 groups

```
SELECT SUBSTR(sourceIP, 1, 7),
SUM(adRevenue)
FROM UserVisits
GROUP BY SUBSTR(sourceIP, 1, 7)
```



[“A Comparison of Approaches to Large-Scale Data Analysis” by A. Pavlo et al., 2004]

# Join Task

## SQL Query:

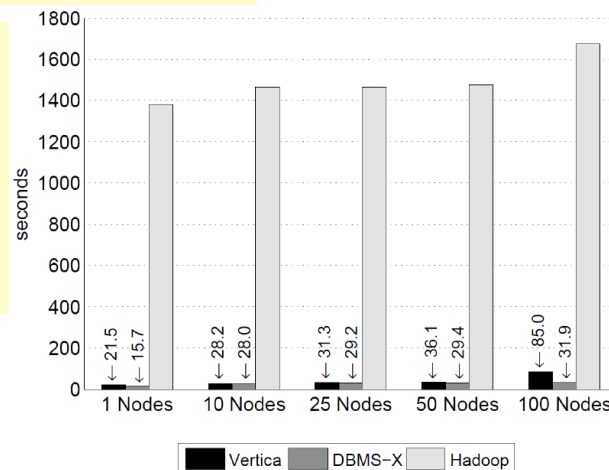
```
SELECT INTO Temp
  UV.sourceIP,
  AVG(R.pageRank) AS avgPageRank,
  SUM(UV.adRevenue) AS totalRevenue
FROM Rankings AS R, UserVisits AS UV
WHERE R.pageURL = UV.destURL
  AND UV.visitDate BETWEEN
    DATE('2000-01-15') AND
    DATE('2000-01-22')
GROUP BY UV.sourceIP
```

```
SELECT sourceIP,
  avgPageRank,
  totalRevenue
FROM Temp
ORDER BY totalRevenue
  DESC LIMIT 1
```

## MapReduce program:

- filter records outside date range, join with rankings file
- compute total ad revenue and average page rank based on source IP
- produce largest total ad revenue record

- Phases run in strict sequential order



[“A Comparison of Approaches to Large-Scale Data Analysis” by A. Pavlo et al., 2004]



# Summary: MapReduce vs Parallel (R)DBMS

- MapReduce: No schema, no index, no high-level language
  - faster loading vs. faster execution
  - easier prototyping vs. easier maintenance
- Fault tolerance
  - restart of single worker vs. restart of transaction
- Installation and tool support
  - easy to setup map/reduce vs. challenging to configure parallel DBMS
  - no tools for tuning vs. tools for automatic performance tuning
- Performance per node
  - results seem to indicate that parallel DBMS achieve same performance as map/reduce in smaller clusters

In a nutshell:

- (R)DBMSs: efficiency, QoS
- MapReduce: cluster scalability




# Spark: improving Hadoop

Credits:


- Matei Zaharia

# Motivation


- MapReduce aiming at “big data” analysis on large, unreliable clusters
  - After initial hype, shortcomings perceived:  
ease of use (programming!), efficiency, tool integration, ...
- ...as soon as organizations started using it widely, users wanted more:
  - More **complex**, multi-stage applications
  - More **interactive** queries
  - More **low-latency** online processing



Iterative job



Interactive mining



Stream processing

# Spark vs Hadoop

- Spark = cluster-computing framework by Berkeley AMPLab
  - Now Apache
- Inherits HDFS, MapReduce from Hadoop
- But:
  - Disk-based comm → in-memory comm
  - Java → Scala

# Resilient Distributed Datasets (RDDs)

- Partitioned collections of records that can be stored **in memory across the cluster**
- Manipulated through a diverse set of transformations
  - *map, filter, join, etc*
- **Fault recovery** without costly replication
  - Remember series of transformations that built RDD (its *lineage*)
  - Can **recompute** lost data based on input files

# Example: Log Mining


- Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")
errors = lines.filter(_.startsWith("ERROR"))
messages = errors.map(_.split('\t')(2))
messages.cache()

messages.filter(_.contains("foo")).count
messages.filter(_.contains("bar")).count
. . .
```



Scala programming language

1 TB data in 5-7 sec (vs 170 sec on disk)



# Ex: Logistic Regression Performance

- Find best line separating two sets of points
- 29 GB dataset
- 20x EC2 m1.xlarge 4-core machines
- Result:



127 s / iteration

first iteration 174 s  
further iterations 6 s

Legend:  
■ Hadoop  
■ Spark

# Conclusion



# Conclusion

- MapReduce = **specialized** (synchronous) distributed processing paradigm
  - Optimized for horizontal scaling in commodity clusters (!), fault tolerance
  - Efficiency? Hardware, energy, ... (see [\[0\]](#), [\[1\]](#), [\[2\]](#), [\[3\]](#) etc.)
    - “Adding more compute servers did not yield significant improvement” [\[src\]](#)
  - Well suited for sets, less so for highly connected data (graphs, arrays)
  - Need to **rewrite algorithms**
- Apache **Hadoop** = MapReduce implementation (HDFS, Java)
- Apache **Spark** = improved MapReduce implementation (HDFS, DSS, Scala)
- **Query languages** on top of MapReduce
  - HLQLs: Pig, Hive, JAQL, ASSET, ...